

CS7643 Deep Learning: Homework 1

Yousef Emam^{1*}

September 26, 2019

^{*1}Y. Emam is with the Institute for Robotics and Intelligent Machines, Georgia Institute of Technology, Atlanta, GA 30332, USA emamy@gatech.edu

1 Gradient Descent

1.1 Optimizer of Unconstrained Opt. Problem

Function to minimize:

$$g(w) = f(w^t) + (w - w^t)^\top \nabla f(w^t) + \frac{\lambda}{2} \|w - w^t\|^2$$

Take the gradient w.r.t. w and set it to 0:

$$\nabla g(w^*) = \nabla f(w^t) + \lambda(w^* - w^t) = 0,$$

which implies:

$$w^* = w^t - \frac{1}{\lambda} \nabla f(w^t).$$

This means that gradient descent is the optimal update with respect to the unconstrained problem if $\eta = \frac{1}{\lambda}$.

1.2 Prove Lemma

Q1.2

INCOMPLETE ANSWER

$$\text{Let } f(v^{(t)}) = \frac{1}{2} \|v^{(t)}\|^2 \Rightarrow \nabla f(v^{(t)}) = v^{(t)}$$

$$\Rightarrow f(w^*) \geq f(v^{(t)}) + \langle w^* - v^{(t)}, \nabla f(v^{(t)}) \rangle$$

$$\Rightarrow \frac{\|w^*\|^2}{2} \geq \frac{1}{2} \|v^{(t)}\|^2 + \langle w^* - v^{(t)}, v^{(t)} \rangle$$

$$\text{Note that } v^{(t)} = \frac{-w^{(t)} + w^{(t-1)}}{\eta}$$

$$\Rightarrow f(w^*) \geq \frac{1}{2} \|v^{(t)}\|^2 + \langle w^* + \frac{w^{(t)} - w^{(t-1)}}{\eta}, v^{(t)} \rangle$$

$$\Rightarrow T f(w^*) \geq \frac{1}{2} \sum_{t=1}^T \|v^{(t)}\|^2 + \sum_{t=1}^T \langle w^* + \frac{w^{(t)} - w^{(t-1)}}{\eta}, v^{(t)} \rangle$$

Figure 1: Question 1b- Scanned Answer

1.3 Bound Convergence of Gradient Descent

Q1.3

$$\bar{w} = \frac{1}{T} \sum_{t=1}^T w^{(t)}$$

$$f(\bar{w}) - f(w^*) = f\left(\frac{1}{T} \sum_{t=1}^T w^{(t)}\right) - f(w^*) \leq \frac{1}{T} \sum_{t=1}^T (f(w^{(t)}) - f(w^*)) \quad (1)$$

$$f(w^*) \geq f(w^{(t)}) + \langle w^* - w^{(t)}, \nabla f(w^{(t)}) \rangle$$

$$\Rightarrow -f(w^*) \leq -f(w^{(t)}) + \langle w^* - w^{(t)}, \nabla f(w^{(t)}) \rangle \quad (2)$$

\Rightarrow Plug (2) into RHS of (1):

$$\begin{aligned} \Rightarrow f(\bar{w}) - f(w^*) &\leq \frac{1}{T} \sum_{t=1}^T \left(\langle w^{(t)} - w^*, \nabla f(w^{(t)}) \rangle \right) \leftarrow \text{Desired Expression} \\ &\leq \frac{1}{T} \left(\frac{\|w^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^T \|\nabla f(w^{(t)})\|^2 \right) \left. \begin{array}{l} \text{Plug in upper bound} \\ \eta = \sqrt{\frac{B^2}{P^2 T}} \end{array} \right\} \\ &\leq \frac{1}{T} \left(\frac{B\sqrt{PT}}{2} + \frac{P\sqrt{T}}{2} \right) \\ &= \boxed{\frac{1}{\sqrt{T}} \left(\frac{B}{2} (B+1) \right)} \end{aligned}$$

Figure 2: Question 1c- Scanned Answer

1.4 SGD Improvement Guarantee

Given function:

$$f(w) = \frac{1}{2}(w-2)^2 + \frac{1}{2}(w+1)^2 = (w - \frac{1}{2})^2 + 2.25$$

Gradient is given by:

$$\nabla f(w) = (w-2) + (w+1)$$

Assume $w^t = 0$, and $N = 2$ is selected:

$$f(0) = (-\frac{1}{2})^2 + 2.25 = 2.5,$$

and

$$w^{t+1} = 0 - \eta \nabla f_2(0) = -\eta.$$

Then:

$$f(w^{t+1}) = (-1/2 - \eta)^2 + 2.25 \geq f(w^t = 0) = (-\frac{1}{2})^2 + 2.25 \forall \eta > 0.$$

The above provides a counter example proving that SGD is not guaranteed to decrease the overall loss function in every iteration.

2 Automatic Differentiation (Figure 7)

2.1 Compute the value of f at $\vec{w} = (1, 2)$

$$f(1, 2) = [7.80207426 * 10^{24}, 2.73105858]$$




Figure 3: Question 2a- Computational Graph

```
def forward(w):  
    w1 = w[0]  
    w2 = w[1]  
    temp = np.exp(w1)+np.exp(2*w2)  
    f1 = np.exp(temp) + np.sin(temp)  
    f2 = w1*w2 + sigmoid(w1)  
    return np.array([f1,f2])
```

Figure 4: Question 2a- Forward Pass

2.2 Compute the Jacobian using num. diff. with $\Delta w = 0.01$

$$\frac{\partial f(1, 2)}{\partial \vec{w}} \approx \begin{bmatrix} 2.18016202 * 10^{25}, 2.19655301 \\ 6.93442399 * 10^{30}, 1.00000000 \end{bmatrix}.$$

Computed using the finite distance formula. Specifically:

$$f'(x) = (f(w + h) - f(w - h)) / 2h.$$

```
# Q2.b
print("Numerical Differentiation: ")
delta_w1 = np.array([0.1, 0])
delta_w2 = np.array([0, 0.1])
print((forward(w+delta_w1)-forward(w-delta_w1))/(2*0.1))
print((forward(w+delta_w2)-forward(w-delta_w2))/(2*0.1))
```

Figure 5: Question 2b

2.3 Compute the Jacobian using forward mode auto-diff.

$$\frac{\partial f(1,2)}{\partial \vec{w}} = \begin{bmatrix} 2.12082367 * 10^{25}, 2.19661193 \\ 8.51957643 * 10^{26}, 1.00000000 \end{bmatrix}.$$

```
def forward_auto(w):
    w1 = w[0]
    w2 = w[1]
    dw1 = 1
    dw2 = 1
    temp = np.exp(w1)+np.exp(2*w2)
    dtemp_dw1 = np.exp(w1)*dw1
    dtemp_dw2 = 2*np.exp(2*w2)*dw2
    f = np.zeros((2,1))
    f[0] = np.exp(temp) + np.sin(temp)
    f[1] = w1*w2 + sigmoid(w1)
    df_dw = np.zeros((2,2))
    df_dw[0,0] = np.exp(temp)*dtemp_dw1 + np.cos(temp)*dtemp_dw1
    df_dw[1,0] = np.exp(temp)*dtemp_dw2 + np.cos(temp)*dtemp_dw2
    df_dw[0,1] = dw1*w2 + sigmoid(w1)*(1-sigmoid(w1))*dw1
    df_dw[1,1] = w1*dw2
    return (f, df_dw)
```

Figure 6: Question 2c- Forward Auto-Diff

2.4 Compute the Jacobian using backward mode auto-diff.

$$\frac{\partial f(1, 2)}{\partial \vec{w}} = \begin{bmatrix} 2.12082367 * 10^{25}, 2.19661193 \\ 8.51957643 * 10^{26}, 1.00000000 \end{bmatrix}.$$

The fact that this is the same result obtained from the forward mode auto-differentiation is not surprising since both forward and backward mode auto-differentiation compute exact derivative.

```
def backward_auto(w):  
    w1 = w[0]  
    w2 = w[1]  
    dw1 = 1  
    dw2 = 1  
    temp = np.exp(w1)+np.exp(2*w2)  
    f = np.zeros((2,1))  
    f[0] = np.exp(temp) + np.sin(temp)  
    f[1] = w1*w2 + sigmoid(w1)  
  
    df1_dtemp = np.exp(temp) + np.cos(temp)  
    df1_w1 = df1_dtemp * np.exp(w1)  
    df1_w2 = df1_dtemp * 2 * np.exp(2*w2)  
    df_dw[0,0] = df1_w1  
    df_dw[1,0] = df1_w2  
    df_dw[0,1] = w2 + sigmoid(w1)*(1-sigmoid(w1))  
    df_dw[1,1] = w1  
    return (f, df_dw)
```

Figure 7: Question 2d- Backward Auto-Diff

2.5 Don't you love that software does this for us?

Yes.

3 Q3: Directed Acyclic Graphs

3.1 If G is a DAG, then G has a topological order

Using the Lemma that every DAG has at least 1 node with in-degree 0, we can construct a topological ordering using the following algorithm:

```
topOrder = {}
while  $G$  is not empty do
    temp  $\leftarrow \{v \in V : \deg_{in}(v) = 0\}$ 
    topOrder  $\leftarrow$  topOrder  $\cup$  enumerate(temp)
     $G \leftarrow$  remove( $G$ ,  $\{v \in V : \deg_{in}(v) = 0\}$ )
end
return topOrder
```

Algorithm 1: Create Topological Order

The loop is guaranteed to terminate since when 0 in-degree nodes ($\{v \in V : \deg_{in}(v) = 0\}$) are removed from the graph, along with their edges, the resulting graph is also a DAG. Therefore, the resulting graph also has at least one node with in-degree 0. Moreover, in this topological ordering, if $n_{v_i} < n_{v_j}$ then there can exist a directed path from v_i to v_j but no path from v_j to v_i .

3.2 If G has a topological order, then G is a DAG

Since many topological orderings can be generated from a DAG, no algorithm can be used to reconstruct the exact DAG given a topological ordering. However, the statement can be proven using the fact that if $n_{v_i} < n_{v_j}$ then there can exist a directed path from v_i to v_j but no path from v_j to v_i . That is simply because if there existed a directed cycle in G , then there must exist two nodes, v_{k_1} and v_{k_2} , such that there is a directed path from v_{k_1} to v_{k_2} and vice-versa. This in turn implies that $n_{v_{k_1}} > n_{v_{k_2}}$ and $n_{v_{k_2}} > n_{v_{k_1}}$ which is impossible.

softmax

September 24, 2019

1 Softmax Classifier

This exercise guides you through the process of classifying images using a Softmax classifier. As part of this you will:

- Implement a fully vectorized loss function for the Softmax classifier
- Calculate the analytical gradient using vectorized code
- Tune hyperparameters on a validation set
- Optimize the loss function with Stochastic Gradient Descent (SGD)
- Visualize the learned weights

```
[80]: # start-up code!
import random

import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
[81]: from load_cifar10_tvt import load_cifar10_train_val

X_train, y_train, X_val, y_val, X_test, y_test = load_cifar10_train_val()
print("Train data shape: ", X_train.shape)
print("Train labels shape: ", y_train.shape)
print("Val data shape: ", X_val.shape)
print("Val labels shape: ", y_val.shape)
```

```
print("Test data shape: ", X_test.shape)
print("Test labels shape: ", y_test.shape)
```

Train, validation and testing sets have been created as

X_i and y_i where $i=\text{train, val, test}$

Train data shape: (3073, 49000)

Train labels shape: (49000,)

Val data shape: (3073, 1000)

Val labels shape: (1000,)

Test data shape: (3073, 1000)

Test labels shape: (1000,)

Code for this section is to be written in `cs231n/classifiers/softmax.py`

```
[156]: # Now, implement the vectorized version in softmax_loss_vectorized.

import time

from cs231n.classifiers.softmax import softmax_loss_vectorized

# gradient check.
from cs231n.gradient_check import grad_check_sparse

W = np.random.randn(10, 3073) * 0.0001

tic = time.time()
loss, grad = softmax_loss_vectorized(W, X_train, y_train, 0.00001)
toc = time.time()
print("vectorized loss: %e computed in %fs" % (loss, toc - tic))

# As a rough sanity check, our loss should be something close to -log(0.1).
print("loss: %f" % loss)
print("sanity check: %f" % (-np.log(0.1)))

f = lambda w: softmax_loss_vectorized(w, X_train, y_train, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

vectorized loss: 2.317796e+00 computed in 0.984168s

loss: 2.317796

sanity check: 2.302585

numerical: 2.060581 analytic: 2.060581, relative error: 8.199045e-09

numerical: -1.910819 analytic: -1.910819, relative error: 4.206250e-09

numerical: 0.747108 analytic: 0.747108, relative error: 2.444263e-08

numerical: 0.396134 analytic: 0.396134, relative error: 5.135847e-09

numerical: 1.727855 analytic: 1.727855, relative error: 5.932070e-08

numerical: 0.781361 analytic: 0.781361, relative error: 8.113437e-09

numerical: -2.858721 analytic: -2.858721, relative error: 8.045210e-09

numerical: 1.416608 analytic: 1.416608, relative error: 2.640299e-08


numerical: -3.852705 analytic: -3.852705, relative error: 1.780070e-09
numerical: 0.527616 analytic: 0.527616, relative error: 1.419871e-08

Code for this section is to be written in `cs231n/classifiers/linear_classifier.py`

```
[157]: # Now that efficient implementations to calculate loss function and gradient of  
↪ the softmax are ready,  
# use it to train the classifier on the cifar-10 data  
# Complete the `train` function in cs231n/classifiers/linear_classifier.py  
  
from cs231n.classifiers.linear_classifier import Softmax  
  
classifier = Softmax()  
loss_hist = classifier.train(  
    X_train,  
    y_train,  
    learning_rate=1e-6,  
    reg=1e-1,  
    num_iters=10000,  
    batch_size=200,  
    verbose=False,  
)  
  
# Training Accuracy  
y_train_pred = classifier.predict(X_train)  
train_accuracy = np.mean(y_train == y_train_pred)  
print("softmax on raw pixels training set accuracy: %f" % (train_accuracy,))  
  
# Plot loss vs. iterations  
plt.plot(loss_hist)  
plt.xlabel("Iteration number")  
plt.ylabel("Loss value")
```

softmax on raw pixels training set accuracy: 0.424755

```
[157]: Text(0, 0.5, 'Loss value')
```



```
[158]: # Complete the `predict` function in cs231n/classifiers/linear_classifier.py
# Evaluate on test set
y_test_pred = classifier.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print("softmax on raw pixels final test set accuracy: %f" % (test_accuracy,))
```

softmax on raw pixels final test set accuracy: 0.380000

```
[159]: # Visualize the learned weights for each class
w = classifier.W[:, :-1] # strip out the bias
w = w.reshape(10, 32, 32, 3)

w_min, w_max = np.min(w), np.max(w)

classes = [
    "plane",
    "car",
    "bird",
    "cat",
```




```

    "deer",
    "dog",
    "frog",
    "horse",
    "ship",
    "truck",
]
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype("uint8"))
    plt.axis("off")
    plt.title(classes[i])

```



two_layer_net

September 24, 2019

1 Implementing a Neural Network

In this exercise we will develop a neural network with fully-connected layers to perform classification, and test it out on the CIFAR-10 dataset.

```
[2]: # A bit of setup

import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The neural network parameters will be stored in a dictionary (model below), where the keys are the parameter names and the values are numpy arrays. Below, we initialize toy data and a toy model that we will use to verify your implementations.

```
[19]: # Create some toy data to check your implementations

input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model():
    model = {}
```

```

    model['W1'] = np.linspace(-0.2, 0.6, num=input_size*hidden_size).
    ↪reshape(input_size, hidden_size)
    model['b1'] = np.linspace(-0.3, 0.7, num=hidden_size)
    model['W2'] = np.linspace(-0.4, 0.1, num=hidden_size*num_classes).
    ↪reshape(hidden_size, num_classes)
    model['b2'] = np.linspace(-0.5, 0.9, num=num_classes)
    return model

def init_toy_data():
    X = np.linspace(-0.2, 0.5, num=num_inputs*input_size).reshape(num_inputs,
    ↪input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

model = init_toy_model()
X, y = init_toy_data()

```

2 Forward pass: compute scores

Open the file `cs231n/classifiers/neural_net.py` and look at the function `two_layer_net`. This function is very similar to the loss functions you have written for the Softmax exercise in HW0: It takes the data and weights and computes the class scores, the loss, and the gradients on the parameters.

Implement the first part of the forward pass which uses the weights and biases to compute the scores for all inputs.

```

[18]: from cs231n.classifiers.neural_net import two_layer_net

scores = two_layer_net(X, model)
print(scores)
correct_scores = [[-0.5328368, 0.20031504, 0.93346689],
                  [-0.59412164, 0.15498488, 0.9040914 ],
                  [-0.67658362, 0.08978957, 0.85616275],
                  [-0.77092643, 0.01339997, 0.79772637],
                  [-0.89110401, -0.08754544, 0.71601312]]

# the difference should be very small. We get 3e-8
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))

```

```

[[-0.5328368  0.20031504  0.93346689]
 [-0.59412164  0.15498488  0.9040914 ]
 [-0.67658362  0.08978957  0.85616275]
 [-0.77092643  0.01339997  0.79772637]
 [-0.89110401 -0.08754544  0.71601312]]

```

Difference between your scores and correct scores:

3.848682278081994e-08

3 Forward pass: compute loss

In the same function, implement the second part that computes the data and regularization loss.

```
[17]: reg = 0.1
      loss, _ = two_layer_net(X, model, y, reg)
      correct_loss = 1.38191946092

      # should be very small, we get 5e-12
      print('Difference between your loss and correct loss:')
      print(np.sum(np.abs(loss - correct_loss)))
```

Difference between your loss and correct loss:
4.6769255135359344e-12

4 Backward pass

Implement the rest of the function. This will compute the gradient of the loss with respect to the variables $W1$, $b1$, $W2$, and $b2$. Now that you (hopefully!) have a correctly implemented forward pass, you can debug your backward pass using a numeric gradient check:

```
[16]: from cs231n.gradient_check import eval_numerical_gradient

      # Use numeric gradient checking to check your implementation of the backward
      # pass.
      # If your implementation is correct, the difference between the numeric and
      # analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

      loss, grads = two_layer_net(X, model, y, reg)

      # these should all be less than 1e-8 or so
      for param_name in grads:
          param_grad_num = eval_numerical_gradient(lambda W: two_layer_net(X, model, y,
          reg)[0], model[param_name], verbose=False)
          print('%s max relative error: %e' % (param_name, rel_error(param_grad_num,
          grads[param_name])))
```

W1 max relative error: 4.426512e-09
b1 max relative error: 5.435432e-08
W2 max relative error: 8.023739e-10
b2 max relative error: 8.190173e-11

5 Train the network

To train the network we will use SGD with Momentum. Last assignment you implemented vanilla SGD. You will now implement the momentum update and the RMSProp update. Open the file `classifier_trainer.py` and familiarize yourself with the `ClassifierTrainer` class. It performs optimization given an arbitrary cost function data, and model. By default it uses vanilla SGD, which we have already implemented for you. First, run the optimization below using Vanilla SGD:

```
[20]: from cs231n.classifier_trainer import ClassifierTrainer

model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient
  ↳Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
                                              reg=0.001,
                                              learning_rate=1e-1, momentum=0.0,
  ↳learning_rate_decay=1,
                                              update='sgd', sample_batches=False,
                                              num_epochs=100,
                                              verbose=False)
print('Final loss with vanilla SGD: %f' % (loss_history[-1], ))
```

```
starting iteration  0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with vanilla SGD: 0.940686
```

Now fill in the **momentum update** in the first missing code block inside the `train` function, and run the same optimization as above but with the momentum update. You should see a much better result in the final obtained loss:

```
[22]: model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient
  ↳Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
```

```

reg=0.001,
learning_rate=1e-1, momentum=0.9,
↪learning_rate_decay=1,
↪sample_batches=False,
update='momentum',
num_epochs=100,
verbose=False)
correct_loss = 0.494394
print('Final loss with momentum SGD: %f. We get: %f' % (loss_history[-1],
↪correct_loss))

```

```

starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with momentum SGD: 0.494394. We get: 0.494394

```

The **RMSProp** update step is given as follows:

```

cache = decay_rate * cache + (1 - decay_rate) * dx**2
x += - learning_rate * dx / np.sqrt(cache + 1e-8)

```

Here, `decay_rate` is a hyperparameter and typical values are [0.9, 0.99, 0.999].

Implement the **RMSProp** update rule inside the `train` function and rerun the optimization:

```

[23]: model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient
↪Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
reg=0.001,
learning_rate=1e-1, momentum=0.9,
↪learning_rate_decay=1,
update='rmsprop',
↪sample_batches=False,
num_epochs=100,
verbose=False)
correct_loss = 0.439368
print('Final loss with RMSProp: %f. We get: %f' % (loss_history[-1],
↪correct_loss))

```

```

starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with RMSProp: 0.429848. We get: 0.439368

```

6 Load the data

Now that you have implemented a two-layer network that passes gradient checks, it's time to load up our favorite CIFAR-10 data so we can use it to train a classifier.

```

[24]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = range(num_training, num_training + num_validation)
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = range(num_training)
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = range(num_test)
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

    # Reshape data to rows
    X_train = X_train.reshape(num_training, -1)

```

```

X_val = X_val.reshape(num_validation, -1)
X_test = X_test.reshape(num_test, -1)

return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

```

```

Train data shape: (49000, 3072)
Train labels shape: (49000,)
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)

```

7 Train a network

To train our network we will use SGD with momentum. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

```

[48]: from cs231n.classifiers.neural_net import init_two_layer_model

model = init_two_layer_model(32*32*3, 70, 10) # input size, hidden size, number
↳ of classes
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train,
↳ X_val, y_val,

                                model, two_layer_net,
                                num_epochs=20, reg=1.0,
                                momentum=0.90, learning_rate_decay
↳ = 0.95,

                                learning_rate=5e-5, verbose=True)

```

```

starting iteration 0
Finished epoch 0 / 20: cost 2.302596, train: 0.085000, val 0.092000, lr
5.000000e-05
starting iteration 10
starting iteration 20
starting iteration 30

```


starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
starting iteration 100
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
starting iteration 460
starting iteration 470
starting iteration 480
Finished epoch 1 / 20: cost 1.703471, train: 0.370000, val 0.359000, lr
4.750000e-05
starting iteration 490

starting iteration 500
starting iteration 510
starting iteration 520
starting iteration 530
starting iteration 540
starting iteration 550
starting iteration 560
starting iteration 570
starting iteration 580
starting iteration 590
starting iteration 600
starting iteration 610
starting iteration 620
starting iteration 630
starting iteration 640
starting iteration 650
starting iteration 660
starting iteration 670
starting iteration 680
starting iteration 690
starting iteration 700
starting iteration 710
starting iteration 720
starting iteration 730
starting iteration 740
starting iteration 750
starting iteration 760
starting iteration 770
starting iteration 780
starting iteration 790
starting iteration 800
starting iteration 810
starting iteration 820
starting iteration 830
starting iteration 840
starting iteration 850
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950
starting iteration 960
starting iteration 970

Finished epoch 2 / 20: cost 1.692839, train: 0.435000, val 0.444000, lr 4.512500e-05

starting iteration 980
starting iteration 990
starting iteration 1000
starting iteration 1010
starting iteration 1020
starting iteration 1030
starting iteration 1040
starting iteration 1050
starting iteration 1060
starting iteration 1070
starting iteration 1080
starting iteration 1090
starting iteration 1100
starting iteration 1110
starting iteration 1120
starting iteration 1130
starting iteration 1140
starting iteration 1150
starting iteration 1160
starting iteration 1170
starting iteration 1180
starting iteration 1190
starting iteration 1200
starting iteration 1210
starting iteration 1220
starting iteration 1230
starting iteration 1240
starting iteration 1250
starting iteration 1260
starting iteration 1270
starting iteration 1280
starting iteration 1290
starting iteration 1300
starting iteration 1310
starting iteration 1320
starting iteration 1330
starting iteration 1340
starting iteration 1350
starting iteration 1360
starting iteration 1370
starting iteration 1380
starting iteration 1390
starting iteration 1400
starting iteration 1410
starting iteration 1420
starting iteration 1430

starting iteration 1440
starting iteration 1450
starting iteration 1460
Finished epoch 3 / 20: cost 1.549167, train: 0.483000, val 0.451000, lr
4.286875e-05
starting iteration 1470
starting iteration 1480
starting iteration 1490
starting iteration 1500
starting iteration 1510
starting iteration 1520
starting iteration 1530
starting iteration 1540
starting iteration 1550
starting iteration 1560
starting iteration 1570
starting iteration 1580
starting iteration 1590
starting iteration 1600
starting iteration 1610
starting iteration 1620
starting iteration 1630
starting iteration 1640
starting iteration 1650
starting iteration 1660
starting iteration 1670
starting iteration 1680
starting iteration 1690
starting iteration 1700
starting iteration 1710
starting iteration 1720
starting iteration 1730
starting iteration 1740
starting iteration 1750
starting iteration 1760
starting iteration 1770
starting iteration 1780
starting iteration 1790
starting iteration 1800
starting iteration 1810
starting iteration 1820
starting iteration 1830
starting iteration 1840
starting iteration 1850
starting iteration 1860
starting iteration 1870
starting iteration 1880
starting iteration 1890

starting iteration 1900
starting iteration 1910
starting iteration 1920
starting iteration 1930
starting iteration 1940
starting iteration 1950
Finished epoch 4 / 20: cost 1.596957, train: 0.476000, val 0.456000, lr
4.072531e-05
starting iteration 1960
starting iteration 1970
starting iteration 1980
starting iteration 1990
starting iteration 2000
starting iteration 2010
starting iteration 2020
starting iteration 2030
starting iteration 2040
starting iteration 2050
starting iteration 2060
starting iteration 2070
starting iteration 2080
starting iteration 2090
starting iteration 2100
starting iteration 2110
starting iteration 2120
starting iteration 2130
starting iteration 2140
starting iteration 2150
starting iteration 2160
starting iteration 2170
starting iteration 2180
starting iteration 2190
starting iteration 2200
starting iteration 2210
starting iteration 2220
starting iteration 2230
starting iteration 2240
starting iteration 2250
starting iteration 2260
starting iteration 2270
starting iteration 2280
starting iteration 2290
starting iteration 2300
starting iteration 2310
starting iteration 2320
starting iteration 2330
starting iteration 2340
starting iteration 2350

starting iteration 2360
starting iteration 2370
starting iteration 2380
starting iteration 2390
starting iteration 2400
starting iteration 2410
starting iteration 2420
starting iteration 2430
starting iteration 2440
Finished epoch 5 / 20: cost 1.591155, train: 0.507000, val 0.465000, lr
3.868905e-05
starting iteration 2450
starting iteration 2460
starting iteration 2470
starting iteration 2480
starting iteration 2490
starting iteration 2500
starting iteration 2510
starting iteration 2520
starting iteration 2530
starting iteration 2540
starting iteration 2550
starting iteration 2560
starting iteration 2570
starting iteration 2580
starting iteration 2590
starting iteration 2600
starting iteration 2610
starting iteration 2620
starting iteration 2630
starting iteration 2640
starting iteration 2650
starting iteration 2660
starting iteration 2670
starting iteration 2680
starting iteration 2690
starting iteration 2700
starting iteration 2710
starting iteration 2720
starting iteration 2730
starting iteration 2740
starting iteration 2750
starting iteration 2760
starting iteration 2770
starting iteration 2780
starting iteration 2790
starting iteration 2800
starting iteration 2810

starting iteration 2820
starting iteration 2830
starting iteration 2840
starting iteration 2850
starting iteration 2860
starting iteration 2870
starting iteration 2880
starting iteration 2890
starting iteration 2900
starting iteration 2910
starting iteration 2920
starting iteration 2930
Finished epoch 6 / 20: cost 1.506842, train: 0.504000, val 0.461000, lr
3.675459e-05
starting iteration 2940
starting iteration 2950
starting iteration 2960
starting iteration 2970
starting iteration 2980
starting iteration 2990
starting iteration 3000
starting iteration 3010
starting iteration 3020
starting iteration 3030
starting iteration 3040
starting iteration 3050
starting iteration 3060
starting iteration 3070
starting iteration 3080
starting iteration 3090
starting iteration 3100
starting iteration 3110
starting iteration 3120
starting iteration 3130
starting iteration 3140
starting iteration 3150
starting iteration 3160
starting iteration 3170
starting iteration 3180
starting iteration 3190
starting iteration 3200
starting iteration 3210
starting iteration 3220
starting iteration 3230
starting iteration 3240
starting iteration 3250
starting iteration 3260
starting iteration 3270

starting iteration 3280
starting iteration 3290
starting iteration 3300
starting iteration 3310
starting iteration 3320
starting iteration 3330
starting iteration 3340
starting iteration 3350
starting iteration 3360
starting iteration 3370
starting iteration 3380
starting iteration 3390
starting iteration 3400
starting iteration 3410
starting iteration 3420
Finished epoch 7 / 20: cost 1.512307, train: 0.542000, val 0.484000, lr
3.491686e-05
starting iteration 3430
starting iteration 3440
starting iteration 3450
starting iteration 3460
starting iteration 3470
starting iteration 3480
starting iteration 3490
starting iteration 3500
starting iteration 3510
starting iteration 3520
starting iteration 3530
starting iteration 3540
starting iteration 3550
starting iteration 3560
starting iteration 3570
starting iteration 3580
starting iteration 3590
starting iteration 3600
starting iteration 3610
starting iteration 3620
starting iteration 3630
starting iteration 3640
starting iteration 3650
starting iteration 3660
starting iteration 3670
starting iteration 3680
starting iteration 3690
starting iteration 3700
starting iteration 3710
starting iteration 3720
starting iteration 3730

starting iteration 3740
starting iteration 3750
starting iteration 3760
starting iteration 3770
starting iteration 3780
starting iteration 3790
starting iteration 3800
starting iteration 3810
starting iteration 3820
starting iteration 3830
starting iteration 3840
starting iteration 3850
starting iteration 3860
starting iteration 3870
starting iteration 3880
starting iteration 3890
starting iteration 3900
starting iteration 3910
Finished epoch 8 / 20: cost 1.628261, train: 0.514000, val 0.492000, lr
3.317102e-05
starting iteration 3920
starting iteration 3930
starting iteration 3940
starting iteration 3950
starting iteration 3960
starting iteration 3970
starting iteration 3980
starting iteration 3990
starting iteration 4000
starting iteration 4010
starting iteration 4020
starting iteration 4030
starting iteration 4040
starting iteration 4050
starting iteration 4060
starting iteration 4070
starting iteration 4080
starting iteration 4090
starting iteration 4100
starting iteration 4110
starting iteration 4120
starting iteration 4130
starting iteration 4140
starting iteration 4150
starting iteration 4160
starting iteration 4170
starting iteration 4180
starting iteration 4190

starting iteration 4200
starting iteration 4210
starting iteration 4220
starting iteration 4230
starting iteration 4240
starting iteration 4250
starting iteration 4260
starting iteration 4270
starting iteration 4280
starting iteration 4290
starting iteration 4300
starting iteration 4310
starting iteration 4320
starting iteration 4330
starting iteration 4340
starting iteration 4350
starting iteration 4360
starting iteration 4370
starting iteration 4380
starting iteration 4390
starting iteration 4400
Finished epoch 9 / 20: cost 1.500103, train: 0.534000, val 0.484000, lr
3.151247e-05
starting iteration 4410
starting iteration 4420
starting iteration 4430
starting iteration 4440
starting iteration 4450
starting iteration 4460
starting iteration 4470
starting iteration 4480
starting iteration 4490
starting iteration 4500
starting iteration 4510
starting iteration 4520
starting iteration 4530
starting iteration 4540
starting iteration 4550
starting iteration 4560
starting iteration 4570
starting iteration 4580
starting iteration 4590
starting iteration 4600
starting iteration 4610
starting iteration 4620
starting iteration 4630
starting iteration 4640
starting iteration 4650

starting iteration 4660
starting iteration 4670
starting iteration 4680
starting iteration 4690
starting iteration 4700
starting iteration 4710
starting iteration 4720
starting iteration 4730
starting iteration 4740
starting iteration 4750
starting iteration 4760
starting iteration 4770
starting iteration 4780
starting iteration 4790
starting iteration 4800
starting iteration 4810
starting iteration 4820
starting iteration 4830
starting iteration 4840
starting iteration 4850
starting iteration 4860
starting iteration 4870
starting iteration 4880
starting iteration 4890
Finished epoch 10 / 20: cost 1.439580, train: 0.553000, val 0.493000, lr
2.993685e-05
starting iteration 4900
starting iteration 4910
starting iteration 4920
starting iteration 4930
starting iteration 4940
starting iteration 4950
starting iteration 4960
starting iteration 4970
starting iteration 4980
starting iteration 4990
starting iteration 5000
starting iteration 5010
starting iteration 5020
starting iteration 5030
starting iteration 5040
starting iteration 5050
starting iteration 5060
starting iteration 5070
starting iteration 5080
starting iteration 5090
starting iteration 5100
starting iteration 5110

starting iteration 5120
starting iteration 5130
starting iteration 5140
starting iteration 5150
starting iteration 5160
starting iteration 5170
starting iteration 5180
starting iteration 5190
starting iteration 5200
starting iteration 5210
starting iteration 5220
starting iteration 5230
starting iteration 5240
starting iteration 5250
starting iteration 5260
starting iteration 5270
starting iteration 5280
starting iteration 5290
starting iteration 5300
starting iteration 5310
starting iteration 5320
starting iteration 5330
starting iteration 5340
starting iteration 5350
starting iteration 5360
starting iteration 5370
starting iteration 5380
Finished epoch 11 / 20: cost 1.420972, train: 0.551000, val 0.485000, lr
2.844000e-05
starting iteration 5390
starting iteration 5400
starting iteration 5410
starting iteration 5420
starting iteration 5430
starting iteration 5440
starting iteration 5450
starting iteration 5460
starting iteration 5470
starting iteration 5480
starting iteration 5490
starting iteration 5500
starting iteration 5510
starting iteration 5520
starting iteration 5530
starting iteration 5540
starting iteration 5550
starting iteration 5560
starting iteration 5570

starting iteration 5580
starting iteration 5590
starting iteration 5600
starting iteration 5610
starting iteration 5620
starting iteration 5630
starting iteration 5640
starting iteration 5650
starting iteration 5660
starting iteration 5670
starting iteration 5680
starting iteration 5690
starting iteration 5700
starting iteration 5710
starting iteration 5720
starting iteration 5730
starting iteration 5740
starting iteration 5750
starting iteration 5760
starting iteration 5770
starting iteration 5780
starting iteration 5790
starting iteration 5800
starting iteration 5810
starting iteration 5820
starting iteration 5830
starting iteration 5840
starting iteration 5850
starting iteration 5860
starting iteration 5870
Finished epoch 12 / 20: cost 1.502760, train: 0.545000, val 0.500000, lr
2.701800e-05
starting iteration 5880
starting iteration 5890
starting iteration 5900
starting iteration 5910
starting iteration 5920
starting iteration 5930
starting iteration 5940
starting iteration 5950
starting iteration 5960
starting iteration 5970
starting iteration 5980
starting iteration 5990
starting iteration 6000
starting iteration 6010
starting iteration 6020
starting iteration 6030

starting iteration 6040
starting iteration 6050
starting iteration 6060
starting iteration 6070
starting iteration 6080
starting iteration 6090
starting iteration 6100
starting iteration 6110
starting iteration 6120
starting iteration 6130
starting iteration 6140
starting iteration 6150
starting iteration 6160
starting iteration 6170
starting iteration 6180
starting iteration 6190
starting iteration 6200
starting iteration 6210
starting iteration 6220
starting iteration 6230
starting iteration 6240
starting iteration 6250
starting iteration 6260
starting iteration 6270
starting iteration 6280
starting iteration 6290
starting iteration 6300
starting iteration 6310
starting iteration 6320
starting iteration 6330
starting iteration 6340
starting iteration 6350
starting iteration 6360
Finished epoch 13 / 20: cost 1.506627, train: 0.535000, val 0.511000, lr
2.566710e-05
starting iteration 6370
starting iteration 6380
starting iteration 6390
starting iteration 6400
starting iteration 6410
starting iteration 6420
starting iteration 6430
starting iteration 6440
starting iteration 6450
starting iteration 6460
starting iteration 6470
starting iteration 6480
starting iteration 6490

```
starting iteration 6500
starting iteration 6510
starting iteration 6520
starting iteration 6530
starting iteration 6540
starting iteration 6550
starting iteration 6560
starting iteration 6570
starting iteration 6580
starting iteration 6590
starting iteration 6600
starting iteration 6610
starting iteration 6620
starting iteration 6630
starting iteration 6640
starting iteration 6650
starting iteration 6660
starting iteration 6670
starting iteration 6680
starting iteration 6690
starting iteration 6700
starting iteration 6710
starting iteration 6720
starting iteration 6730
starting iteration 6740
starting iteration 6750
starting iteration 6760
starting iteration 6770
starting iteration 6780
starting iteration 6790
starting iteration 6800
starting iteration 6810
starting iteration 6820
starting iteration 6830
starting iteration 6840
starting iteration 6850
Finished epoch 14 / 20: cost 1.482915, train: 0.572000, val 0.489000, lr
2.438375e-05
starting iteration 6860
starting iteration 6870
starting iteration 6880
starting iteration 6890
starting iteration 6900
starting iteration 6910
starting iteration 6920
starting iteration 6930
starting iteration 6940
starting iteration 6950
```

starting iteration 6960
starting iteration 6970
starting iteration 6980
starting iteration 6990
starting iteration 7000
starting iteration 7010
starting iteration 7020
starting iteration 7030
starting iteration 7040
starting iteration 7050
starting iteration 7060
starting iteration 7070
starting iteration 7080
starting iteration 7090
starting iteration 7100
starting iteration 7110
starting iteration 7120
starting iteration 7130
starting iteration 7140
starting iteration 7150
starting iteration 7160
starting iteration 7170
starting iteration 7180
starting iteration 7190
starting iteration 7200
starting iteration 7210
starting iteration 7220
starting iteration 7230
starting iteration 7240
starting iteration 7250
starting iteration 7260
starting iteration 7270
starting iteration 7280
starting iteration 7290
starting iteration 7300
starting iteration 7310
starting iteration 7320
starting iteration 7330
starting iteration 7340
Finished epoch 15 / 20: cost 1.639811, train: 0.568000, val 0.502000, lr
2.316456e-05
starting iteration 7350
starting iteration 7360
starting iteration 7370
starting iteration 7380
starting iteration 7390
starting iteration 7400
starting iteration 7410

starting iteration 7420
starting iteration 7430
starting iteration 7440
starting iteration 7450
starting iteration 7460
starting iteration 7470
starting iteration 7480
starting iteration 7490
starting iteration 7500
starting iteration 7510
starting iteration 7520
starting iteration 7530
starting iteration 7540
starting iteration 7550
starting iteration 7560
starting iteration 7570
starting iteration 7580
starting iteration 7590
starting iteration 7600
starting iteration 7610
starting iteration 7620
starting iteration 7630
starting iteration 7640
starting iteration 7650
starting iteration 7660
starting iteration 7670
starting iteration 7680
starting iteration 7690
starting iteration 7700
starting iteration 7710
starting iteration 7720
starting iteration 7730
starting iteration 7740
starting iteration 7750
starting iteration 7760
starting iteration 7770
starting iteration 7780
starting iteration 7790
starting iteration 7800
starting iteration 7810
starting iteration 7820
starting iteration 7830
Finished epoch 16 / 20: cost 1.512649, train: 0.550000, val 0.506000, lr
2.200633e-05
starting iteration 7840
starting iteration 7850
starting iteration 7860
starting iteration 7870

starting iteration 7880
starting iteration 7890
starting iteration 7900
starting iteration 7910
starting iteration 7920
starting iteration 7930
starting iteration 7940
starting iteration 7950
starting iteration 7960
starting iteration 7970
starting iteration 7980
starting iteration 7990
starting iteration 8000
starting iteration 8010
starting iteration 8020
starting iteration 8030
starting iteration 8040
starting iteration 8050
starting iteration 8060
starting iteration 8070
starting iteration 8080
starting iteration 8090
starting iteration 8100
starting iteration 8110
starting iteration 8120
starting iteration 8130
starting iteration 8140
starting iteration 8150
starting iteration 8160
starting iteration 8170
starting iteration 8180
starting iteration 8190
starting iteration 8200
starting iteration 8210
starting iteration 8220
starting iteration 8230
starting iteration 8240
starting iteration 8250
starting iteration 8260
starting iteration 8270
starting iteration 8280
starting iteration 8290
starting iteration 8300
starting iteration 8310
starting iteration 8320
Finished epoch 17 / 20: cost 1.473777, train: 0.555000, val 0.517000, lr
2.090602e-05
starting iteration 8330

starting iteration 8340
starting iteration 8350
starting iteration 8360
starting iteration 8370
starting iteration 8380
starting iteration 8390
starting iteration 8400
starting iteration 8410
starting iteration 8420
starting iteration 8430
starting iteration 8440
starting iteration 8450
starting iteration 8460
starting iteration 8470
starting iteration 8480
starting iteration 8490
starting iteration 8500
starting iteration 8510
starting iteration 8520
starting iteration 8530
starting iteration 8540
starting iteration 8550
starting iteration 8560
starting iteration 8570
starting iteration 8580
starting iteration 8590
starting iteration 8600
starting iteration 8610
starting iteration 8620
starting iteration 8630
starting iteration 8640
starting iteration 8650
starting iteration 8660
starting iteration 8670
starting iteration 8680
starting iteration 8690
starting iteration 8700
starting iteration 8710
starting iteration 8720
starting iteration 8730
starting iteration 8740
starting iteration 8750
starting iteration 8760
starting iteration 8770
starting iteration 8780
starting iteration 8790
starting iteration 8800
starting iteration 8810

Finished epoch 18 / 20: cost 1.556926, train: 0.562000, val 0.509000, lr 1.986072e-05

starting iteration 8820
starting iteration 8830
starting iteration 8840
starting iteration 8850
starting iteration 8860
starting iteration 8870
starting iteration 8880
starting iteration 8890
starting iteration 8900
starting iteration 8910
starting iteration 8920
starting iteration 8930
starting iteration 8940
starting iteration 8950
starting iteration 8960
starting iteration 8970
starting iteration 8980
starting iteration 8990
starting iteration 9000
starting iteration 9010
starting iteration 9020
starting iteration 9030
starting iteration 9040
starting iteration 9050
starting iteration 9060
starting iteration 9070
starting iteration 9080
starting iteration 9090
starting iteration 9100
starting iteration 9110
starting iteration 9120
starting iteration 9130
starting iteration 9140
starting iteration 9150
starting iteration 9160
starting iteration 9170
starting iteration 9180
starting iteration 9190
starting iteration 9200
starting iteration 9210
starting iteration 9220
starting iteration 9230
starting iteration 9240
starting iteration 9250
starting iteration 9260
starting iteration 9270

starting iteration 9280
starting iteration 9290
starting iteration 9300
Finished epoch 19 / 20: cost 1.502153, train: 0.546000, val 0.513000, lr
1.886768e-05
starting iteration 9310
starting iteration 9320
starting iteration 9330
starting iteration 9340
starting iteration 9350
starting iteration 9360
starting iteration 9370
starting iteration 9380
starting iteration 9390
starting iteration 9400
starting iteration 9410
starting iteration 9420
starting iteration 9430
starting iteration 9440
starting iteration 9450
starting iteration 9460
starting iteration 9470
starting iteration 9480
starting iteration 9490
starting iteration 9500
starting iteration 9510
starting iteration 9520
starting iteration 9530
starting iteration 9540
starting iteration 9550
starting iteration 9560
starting iteration 9570
starting iteration 9580
starting iteration 9590
starting iteration 9600
starting iteration 9610
starting iteration 9620
starting iteration 9630
starting iteration 9640
starting iteration 9650
starting iteration 9660
starting iteration 9670
starting iteration 9680
starting iteration 9690
starting iteration 9700
starting iteration 9710
starting iteration 9720
starting iteration 9730

```
starting iteration 9740
starting iteration 9750
starting iteration 9760
starting iteration 9770
starting iteration 9780
starting iteration 9790
Finished epoch 20 / 20: cost 1.272936, train: 0.561000, val 0.521000, lr
1.792430e-05
finished optimization. best validation accuracy: 0.521000
```

8 Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.37 on the validation set. This isn't very good.


One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
[49]: # Plot the loss function and train / validation accuracies
plt.subplot(2, 1, 1)
plt.plot(loss_history)
plt.title('Loss history')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
plt.plot(train_acc)
plt.plot(val_acc)
plt.legend(['Training accuracy', 'Validation accuracy'], loc='lower right')
plt.xlabel('Epoch')
plt.ylabel('Clasification accuracy')
```

```
[49]: Text(0, 0.5, 'Clasification accuracy')
```




```
[50]: from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(model):
    plt.imshow(visualize_grid(model['W1'].T.reshape(-1, 32, 32, 3), padding=3).
        ↪astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```



9 Tune your hyperparameters

What's wrong?. Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

Tuning. Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the momentum and learning rate decay parameters, but you should be able to get good performance using the default values.

Approximate results. You should be aim to achieve a classification accuracy of greater than 50% on the validation set. Our best network gets over 56% on the validation set.

Experiment: Your goal in this exercise is to get as good of a result on CIFAR-10 as you can, with a fully-connected Neural Network. For every 1% above 56% on the Test set we will award you with one extra bonus point. Feel free to implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

```
[52]: best_model = None # store the best model into this
```

```
#####
```



```

# TODO: Tune hyperparameters using the validation set. Store your best trained
↪#
# model in best_model.
↪#
#
↪#
# To help debug your network, it may help to use visualizations similar to the
↪#
# ones we used above; these visualizations will have significant qualitative
↪#
# differences from the ones we saw above for the poorly tuned network.
↪#
#
↪#
# Tweaking hyperparameters by hand can be fun, but you might find it useful to
↪#
# write code to sweep through possible combinations of hyperparameters
↪#
# automatically like we did on the previous assignment.
↪#
#####
# input size, hidden size, number of classes
model = init_two_layer_model(32*32*3, 1000, 10)
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train,
↪X_val, y_val,
                                model, two_layer_net,
                                num_epochs=20, reg=1.0,
                                momentum=0.90, learning_rate_decay
↪= 0.95,
                                learning_rate=5e-5, verbose=True)
#####
#                               END OF YOUR CODE
↪#
#####

```

```

starting iteration 0
Finished epoch 0 / 20: cost 2.302740, train: 0.088000, val 0.085000, lr
5.000000e-05
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70

```

starting iteration 80
starting iteration 90
starting iteration 100
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
starting iteration 460
starting iteration 470
starting iteration 480
Finished epoch 1 / 20: cost 1.800183, train: 0.359000, val 0.386000, lr
4.750000e-05
starting iteration 490
starting iteration 500
starting iteration 510
starting iteration 520
starting iteration 530

starting iteration 540
starting iteration 550
starting iteration 560
starting iteration 570
starting iteration 580
starting iteration 590
starting iteration 600
starting iteration 610
starting iteration 620
starting iteration 630
starting iteration 640
starting iteration 650
starting iteration 660
starting iteration 670
starting iteration 680
starting iteration 690
starting iteration 700
starting iteration 710
starting iteration 720
starting iteration 730
starting iteration 740
starting iteration 750
starting iteration 760
starting iteration 770
starting iteration 780
starting iteration 790
starting iteration 800
starting iteration 810
starting iteration 820
starting iteration 830
starting iteration 840
starting iteration 850
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950
starting iteration 960
starting iteration 970
Finished epoch 2 / 20: cost 1.586629, train: 0.431000, val 0.452000, lr
4.512500e-05
starting iteration 980
starting iteration 990

starting iteration 1000
starting iteration 1010
starting iteration 1020
starting iteration 1030
starting iteration 1040
starting iteration 1050
starting iteration 1060
starting iteration 1070
starting iteration 1080
starting iteration 1090
starting iteration 1100
starting iteration 1110
starting iteration 1120
starting iteration 1130
starting iteration 1140
starting iteration 1150
starting iteration 1160
starting iteration 1170
starting iteration 1180
starting iteration 1190
starting iteration 1200
starting iteration 1210
starting iteration 1220
starting iteration 1230
starting iteration 1240
starting iteration 1250
starting iteration 1260
starting iteration 1270
starting iteration 1280
starting iteration 1290
starting iteration 1300
starting iteration 1310
starting iteration 1320
starting iteration 1330
starting iteration 1340
starting iteration 1350
starting iteration 1360
starting iteration 1370
starting iteration 1380
starting iteration 1390
starting iteration 1400
starting iteration 1410
starting iteration 1420
starting iteration 1430
starting iteration 1440
starting iteration 1450
starting iteration 1460
Finished epoch 3 / 20: cost 1.681464, train: 0.477000, val 0.466000, lr

4.286875e-05

starting iteration 1470
starting iteration 1480
starting iteration 1490
starting iteration 1500
starting iteration 1510
starting iteration 1520
starting iteration 1530
starting iteration 1540
starting iteration 1550
starting iteration 1560
starting iteration 1570
starting iteration 1580
starting iteration 1590
starting iteration 1600
starting iteration 1610
starting iteration 1620
starting iteration 1630
starting iteration 1640
starting iteration 1650
starting iteration 1660
starting iteration 1670
starting iteration 1680
starting iteration 1690
starting iteration 1700
starting iteration 1710
starting iteration 1720
starting iteration 1730
starting iteration 1740
starting iteration 1750
starting iteration 1760
starting iteration 1770
starting iteration 1780
starting iteration 1790
starting iteration 1800
starting iteration 1810
starting iteration 1820
starting iteration 1830
starting iteration 1840
starting iteration 1850
starting iteration 1860
starting iteration 1870
starting iteration 1880
starting iteration 1890
starting iteration 1900
starting iteration 1910
starting iteration 1920
starting iteration 1930

starting iteration 1940
starting iteration 1950
Finished epoch 4 / 20: cost 1.542658, train: 0.497000, val 0.479000, lr
4.072531e-05
starting iteration 1960
starting iteration 1970
starting iteration 1980
starting iteration 1990
starting iteration 2000
starting iteration 2010
starting iteration 2020
starting iteration 2030
starting iteration 2040
starting iteration 2050
starting iteration 2060
starting iteration 2070
starting iteration 2080
starting iteration 2090
starting iteration 2100
starting iteration 2110
starting iteration 2120
starting iteration 2130
starting iteration 2140
starting iteration 2150
starting iteration 2160
starting iteration 2170
starting iteration 2180
starting iteration 2190
starting iteration 2200
starting iteration 2210
starting iteration 2220
starting iteration 2230
starting iteration 2240
starting iteration 2250
starting iteration 2260
starting iteration 2270
starting iteration 2280
starting iteration 2290
starting iteration 2300
starting iteration 2310
starting iteration 2320
starting iteration 2330
starting iteration 2340
starting iteration 2350
starting iteration 2360
starting iteration 2370
starting iteration 2380
starting iteration 2390

starting iteration 2400
starting iteration 2410
starting iteration 2420
starting iteration 2430
starting iteration 2440
Finished epoch 5 / 20: cost 1.421326, train: 0.503000, val 0.498000, lr
3.868905e-05
starting iteration 2450
starting iteration 2460
starting iteration 2470
starting iteration 2480
starting iteration 2490
starting iteration 2500
starting iteration 2510
starting iteration 2520
starting iteration 2530
starting iteration 2540
starting iteration 2550
starting iteration 2560
starting iteration 2570
starting iteration 2580
starting iteration 2590
starting iteration 2600
starting iteration 2610
starting iteration 2620
starting iteration 2630
starting iteration 2640
starting iteration 2650
starting iteration 2660
starting iteration 2670
starting iteration 2680
starting iteration 2690
starting iteration 2700
starting iteration 2710
starting iteration 2720
starting iteration 2730
starting iteration 2740
starting iteration 2750
starting iteration 2760
starting iteration 2770
starting iteration 2780
starting iteration 2790
starting iteration 2800
starting iteration 2810
starting iteration 2820
starting iteration 2830
starting iteration 2840
starting iteration 2850

starting iteration 2860
starting iteration 2870
starting iteration 2880
starting iteration 2890
starting iteration 2900
starting iteration 2910
starting iteration 2920
starting iteration 2930
Finished epoch 6 / 20: cost 1.476008, train: 0.514000, val 0.500000, lr
3.675459e-05
starting iteration 2940
starting iteration 2950
starting iteration 2960
starting iteration 2970
starting iteration 2980
starting iteration 2990
starting iteration 3000
starting iteration 3010
starting iteration 3020
starting iteration 3030
starting iteration 3040
starting iteration 3050
starting iteration 3060
starting iteration 3070
starting iteration 3080
starting iteration 3090
starting iteration 3100
starting iteration 3110
starting iteration 3120
starting iteration 3130
starting iteration 3140
starting iteration 3150
starting iteration 3160
starting iteration 3170
starting iteration 3180
starting iteration 3190
starting iteration 3200
starting iteration 3210
starting iteration 3220
starting iteration 3230
starting iteration 3240
starting iteration 3250
starting iteration 3260
starting iteration 3270
starting iteration 3280
starting iteration 3290
starting iteration 3300
starting iteration 3310

starting iteration 3320
starting iteration 3330
starting iteration 3340
starting iteration 3350
starting iteration 3360
starting iteration 3370
starting iteration 3380
starting iteration 3390
starting iteration 3400
starting iteration 3410
starting iteration 3420
Finished epoch 7 / 20: cost 1.537347, train: 0.551000, val 0.528000, lr
3.491686e-05
starting iteration 3430
starting iteration 3440
starting iteration 3450
starting iteration 3460
starting iteration 3470
starting iteration 3480
starting iteration 3490
starting iteration 3500
starting iteration 3510
starting iteration 3520
starting iteration 3530
starting iteration 3540
starting iteration 3550
starting iteration 3560
starting iteration 3570
starting iteration 3580
starting iteration 3590
starting iteration 3600
starting iteration 3610
starting iteration 3620
starting iteration 3630
starting iteration 3640
starting iteration 3650
starting iteration 3660
starting iteration 3670
starting iteration 3680
starting iteration 3690
starting iteration 3700
starting iteration 3710
starting iteration 3720
starting iteration 3730
starting iteration 3740
starting iteration 3750
starting iteration 3760
starting iteration 3770

starting iteration 3780
starting iteration 3790
starting iteration 3800
starting iteration 3810
starting iteration 3820
starting iteration 3830
starting iteration 3840
starting iteration 3850
starting iteration 3860
starting iteration 3870
starting iteration 3880
starting iteration 3890
starting iteration 3900
starting iteration 3910
Finished epoch 8 / 20: cost 1.493411, train: 0.535000, val 0.482000, lr
3.317102e-05
starting iteration 3920
starting iteration 3930
starting iteration 3940
starting iteration 3950
starting iteration 3960
starting iteration 3970
starting iteration 3980
starting iteration 3990
starting iteration 4000
starting iteration 4010
starting iteration 4020
starting iteration 4030
starting iteration 4040
starting iteration 4050
starting iteration 4060
starting iteration 4070
starting iteration 4080
starting iteration 4090
starting iteration 4100
starting iteration 4110
starting iteration 4120
starting iteration 4130
starting iteration 4140
starting iteration 4150
starting iteration 4160
starting iteration 4170
starting iteration 4180
starting iteration 4190
starting iteration 4200
starting iteration 4210
starting iteration 4220
starting iteration 4230

starting iteration 4240
starting iteration 4250
starting iteration 4260
starting iteration 4270
starting iteration 4280
starting iteration 4290
starting iteration 4300
starting iteration 4310
starting iteration 4320
starting iteration 4330
starting iteration 4340
starting iteration 4350
starting iteration 4360
starting iteration 4370
starting iteration 4380
starting iteration 4390
starting iteration 4400
Finished epoch 9 / 20: cost 1.323490, train: 0.548000, val 0.510000, lr
3.151247e-05
starting iteration 4410
starting iteration 4420
starting iteration 4430
starting iteration 4440
starting iteration 4450
starting iteration 4460
starting iteration 4470
starting iteration 4480
starting iteration 4490
starting iteration 4500
starting iteration 4510
starting iteration 4520
starting iteration 4530
starting iteration 4540
starting iteration 4550
starting iteration 4560
starting iteration 4570
starting iteration 4580
starting iteration 4590
starting iteration 4600
starting iteration 4610
starting iteration 4620
starting iteration 4630
starting iteration 4640
starting iteration 4650
starting iteration 4660
starting iteration 4670
starting iteration 4680
starting iteration 4690

```
starting iteration 4700
starting iteration 4710
starting iteration 4720
starting iteration 4730
starting iteration 4740
starting iteration 4750
starting iteration 4760
starting iteration 4770
starting iteration 4780
starting iteration 4790
starting iteration 4800
starting iteration 4810
starting iteration 4820
starting iteration 4830
starting iteration 4840
starting iteration 4850
starting iteration 4860
starting iteration 4870
starting iteration 4880
starting iteration 4890
Finished epoch 10 / 20: cost 1.398117, train: 0.565000, val 0.520000, lr
2.993685e-05
starting iteration 4900
starting iteration 4910
starting iteration 4920
starting iteration 4930
starting iteration 4940
starting iteration 4950
starting iteration 4960
starting iteration 4970
starting iteration 4980
starting iteration 4990
starting iteration 5000
starting iteration 5010
starting iteration 5020
starting iteration 5030
starting iteration 5040
starting iteration 5050
starting iteration 5060
starting iteration 5070
starting iteration 5080
starting iteration 5090
starting iteration 5100
starting iteration 5110
starting iteration 5120
starting iteration 5130
starting iteration 5140
starting iteration 5150
```

starting iteration 5160
starting iteration 5170
starting iteration 5180
starting iteration 5190
starting iteration 5200
starting iteration 5210
starting iteration 5220
starting iteration 5230
starting iteration 5240
starting iteration 5250
starting iteration 5260
starting iteration 5270
starting iteration 5280
starting iteration 5290
starting iteration 5300
starting iteration 5310
starting iteration 5320
starting iteration 5330
starting iteration 5340
starting iteration 5350
starting iteration 5360
starting iteration 5370
starting iteration 5380
Finished epoch 11 / 20: cost 1.619243, train: 0.569000, val 0.508000, lr
2.844000e-05
starting iteration 5390
starting iteration 5400
starting iteration 5410
starting iteration 5420
starting iteration 5430
starting iteration 5440
starting iteration 5450
starting iteration 5460
starting iteration 5470
starting iteration 5480
starting iteration 5490
starting iteration 5500
starting iteration 5510
starting iteration 5520
starting iteration 5530
starting iteration 5540
starting iteration 5550
starting iteration 5560
starting iteration 5570
starting iteration 5580
starting iteration 5590
starting iteration 5600
starting iteration 5610

starting iteration 5620
starting iteration 5630
starting iteration 5640
starting iteration 5650
starting iteration 5660
starting iteration 5670
starting iteration 5680
starting iteration 5690
starting iteration 5700
starting iteration 5710
starting iteration 5720
starting iteration 5730
starting iteration 5740
starting iteration 5750
starting iteration 5760
starting iteration 5770
starting iteration 5780
starting iteration 5790
starting iteration 5800
starting iteration 5810
starting iteration 5820
starting iteration 5830
starting iteration 5840
starting iteration 5850
starting iteration 5860
starting iteration 5870
Finished epoch 12 / 20: cost 1.483477, train: 0.563000, val 0.523000, lr
2.701800e-05
starting iteration 5880
starting iteration 5890
starting iteration 5900
starting iteration 5910
starting iteration 5920
starting iteration 5930
starting iteration 5940
starting iteration 5950
starting iteration 5960
starting iteration 5970
starting iteration 5980
starting iteration 5990
starting iteration 6000
starting iteration 6010
starting iteration 6020
starting iteration 6030
starting iteration 6040
starting iteration 6050
starting iteration 6060
starting iteration 6070

starting iteration 6080
starting iteration 6090
starting iteration 6100
starting iteration 6110
starting iteration 6120
starting iteration 6130
starting iteration 6140
starting iteration 6150
starting iteration 6160
starting iteration 6170
starting iteration 6180
starting iteration 6190
starting iteration 6200
starting iteration 6210
starting iteration 6220
starting iteration 6230
starting iteration 6240
starting iteration 6250
starting iteration 6260
starting iteration 6270
starting iteration 6280
starting iteration 6290
starting iteration 6300
starting iteration 6310
starting iteration 6320
starting iteration 6330
starting iteration 6340
starting iteration 6350
starting iteration 6360
Finished epoch 13 / 20: cost 1.403796, train: 0.542000, val 0.523000, lr
2.566710e-05
starting iteration 6370
starting iteration 6380
starting iteration 6390
starting iteration 6400
starting iteration 6410
starting iteration 6420
starting iteration 6430
starting iteration 6440
starting iteration 6450
starting iteration 6460
starting iteration 6470
starting iteration 6480
starting iteration 6490
starting iteration 6500
starting iteration 6510
starting iteration 6520
starting iteration 6530

```
starting iteration 6540
starting iteration 6550
starting iteration 6560
starting iteration 6570
starting iteration 6580
starting iteration 6590
starting iteration 6600
starting iteration 6610
starting iteration 6620
starting iteration 6630
starting iteration 6640
starting iteration 6650
starting iteration 6660
starting iteration 6670
starting iteration 6680
starting iteration 6690
starting iteration 6700
starting iteration 6710
starting iteration 6720
starting iteration 6730
starting iteration 6740
starting iteration 6750
starting iteration 6760
starting iteration 6770
starting iteration 6780
starting iteration 6790
starting iteration 6800
starting iteration 6810
starting iteration 6820
starting iteration 6830
starting iteration 6840
starting iteration 6850
Finished epoch 14 / 20: cost 1.512430, train: 0.603000, val 0.525000, lr
2.438375e-05
starting iteration 6860
starting iteration 6870
starting iteration 6880
starting iteration 6890
starting iteration 6900
starting iteration 6910
starting iteration 6920
starting iteration 6930
starting iteration 6940
starting iteration 6950
starting iteration 6960
starting iteration 6970
starting iteration 6980
starting iteration 6990
```



```
starting iteration 7000
starting iteration 7010
starting iteration 7020
starting iteration 7030
starting iteration 7040
starting iteration 7050
starting iteration 7060
starting iteration 7070
starting iteration 7080
starting iteration 7090
starting iteration 7100
starting iteration 7110
starting iteration 7120
starting iteration 7130
starting iteration 7140
starting iteration 7150
starting iteration 7160
starting iteration 7170
starting iteration 7180
starting iteration 7190
starting iteration 7200
starting iteration 7210
starting iteration 7220
starting iteration 7230
starting iteration 7240
starting iteration 7250
starting iteration 7260
starting iteration 7270
starting iteration 7280
starting iteration 7290
starting iteration 7300
starting iteration 7310
starting iteration 7320
starting iteration 7330
starting iteration 7340
Finished epoch 15 / 20: cost 1.381355, train: 0.591000, val 0.541000, lr
2.316456e-05
starting iteration 7350
starting iteration 7360
starting iteration 7370
starting iteration 7380
starting iteration 7390
starting iteration 7400
starting iteration 7410
starting iteration 7420
starting iteration 7430
starting iteration 7440
starting iteration 7450
```

starting iteration 7460
starting iteration 7470
starting iteration 7480
starting iteration 7490
starting iteration 7500
starting iteration 7510
starting iteration 7520
starting iteration 7530
starting iteration 7540
starting iteration 7550
starting iteration 7560
starting iteration 7570
starting iteration 7580
starting iteration 7590
starting iteration 7600
starting iteration 7610
starting iteration 7620
starting iteration 7630
starting iteration 7640
starting iteration 7650
starting iteration 7660
starting iteration 7670
starting iteration 7680
starting iteration 7690
starting iteration 7700
starting iteration 7710
starting iteration 7720
starting iteration 7730
starting iteration 7740
starting iteration 7750
starting iteration 7760
starting iteration 7770
starting iteration 7780
starting iteration 7790
starting iteration 7800
starting iteration 7810
starting iteration 7820
starting iteration 7830
Finished epoch 16 / 20: cost 1.410587, train: 0.602000, val 0.540000, lr
2.200633e-05
starting iteration 7840
starting iteration 7850
starting iteration 7860
starting iteration 7870
starting iteration 7880
starting iteration 7890
starting iteration 7900
starting iteration 7910

starting iteration 7920
starting iteration 7930
starting iteration 7940
starting iteration 7950
starting iteration 7960
starting iteration 7970
starting iteration 7980
starting iteration 7990
starting iteration 8000
starting iteration 8010
starting iteration 8020
starting iteration 8030
starting iteration 8040
starting iteration 8050
starting iteration 8060
starting iteration 8070
starting iteration 8080
starting iteration 8090
starting iteration 8100
starting iteration 8110
starting iteration 8120
starting iteration 8130
starting iteration 8140
starting iteration 8150
starting iteration 8160
starting iteration 8170
starting iteration 8180
starting iteration 8190
starting iteration 8200
starting iteration 8210
starting iteration 8220
starting iteration 8230
starting iteration 8240
starting iteration 8250
starting iteration 8260
starting iteration 8270
starting iteration 8280
starting iteration 8290
starting iteration 8300
starting iteration 8310
starting iteration 8320
Finished epoch 17 / 20: cost 1.482933, train: 0.619000, val 0.549000, lr
2.090602e-05
starting iteration 8330
starting iteration 8340
starting iteration 8350
starting iteration 8360
starting iteration 8370

starting iteration 8380
starting iteration 8390
starting iteration 8400
starting iteration 8410
starting iteration 8420
starting iteration 8430
starting iteration 8440
starting iteration 8450
starting iteration 8460
starting iteration 8470
starting iteration 8480
starting iteration 8490
starting iteration 8500
starting iteration 8510
starting iteration 8520
starting iteration 8530
starting iteration 8540
starting iteration 8550
starting iteration 8560
starting iteration 8570
starting iteration 8580
starting iteration 8590
starting iteration 8600
starting iteration 8610
starting iteration 8620
starting iteration 8630
starting iteration 8640
starting iteration 8650
starting iteration 8660
starting iteration 8670
starting iteration 8680
starting iteration 8690
starting iteration 8700
starting iteration 8710
starting iteration 8720
starting iteration 8730
starting iteration 8740
starting iteration 8750
starting iteration 8760
starting iteration 8770
starting iteration 8780
starting iteration 8790
starting iteration 8800
starting iteration 8810
Finished epoch 18 / 20: cost 1.184236, train: 0.616000, val 0.528000, lr
1.986072e-05
starting iteration 8820
starting iteration 8830


starting iteration 8840
starting iteration 8850
starting iteration 8860
starting iteration 8870
starting iteration 8880
starting iteration 8890
starting iteration 8900
starting iteration 8910
starting iteration 8920
starting iteration 8930
starting iteration 8940
starting iteration 8950
starting iteration 8960
starting iteration 8970
starting iteration 8980
starting iteration 8990
starting iteration 9000
starting iteration 9010
starting iteration 9020
starting iteration 9030
starting iteration 9040
starting iteration 9050
starting iteration 9060
starting iteration 9070
starting iteration 9080
starting iteration 9090
starting iteration 9100
starting iteration 9110
starting iteration 9120
starting iteration 9130
starting iteration 9140
starting iteration 9150
starting iteration 9160
starting iteration 9170
starting iteration 9180
starting iteration 9190
starting iteration 9200
starting iteration 9210
starting iteration 9220
starting iteration 9230
starting iteration 9240
starting iteration 9250
starting iteration 9260
starting iteration 9270
starting iteration 9280
starting iteration 9290
starting iteration 9300
Finished epoch 19 / 20: cost 1.375873, train: 0.629000, val 0.545000, lr

1.886768e-05

starting iteration 9310
starting iteration 9320
starting iteration 9330
starting iteration 9340
starting iteration 9350
starting iteration 9360
starting iteration 9370
starting iteration 9380
starting iteration 9390
starting iteration 9400
starting iteration 9410
starting iteration 9420
starting iteration 9430
starting iteration 9440
starting iteration 9450
starting iteration 9460
starting iteration 9470
starting iteration 9480
starting iteration 9490
starting iteration 9500
starting iteration 9510
starting iteration 9520
starting iteration 9530
starting iteration 9540
starting iteration 9550
starting iteration 9560
starting iteration 9570
starting iteration 9580
starting iteration 9590
starting iteration 9600
starting iteration 9610
starting iteration 9620
starting iteration 9630
starting iteration 9640
starting iteration 9650
starting iteration 9660
starting iteration 9670
starting iteration 9680
starting iteration 9690
starting iteration 9700
starting iteration 9710
starting iteration 9720
starting iteration 9730
starting iteration 9740
starting iteration 9750
starting iteration 9760
starting iteration 9770

```
starting iteration 9780
starting iteration 9790
Finished epoch 20 / 20: cost 1.288014, train: 0.616000, val 0.535000, lr
1.792430e-05
finished optimization. best validation accuracy: 0.549000
```

```
[53]: # visualize the weights
show_net_weights(best_model)
```



10 Run on the test set

When you are done experimenting, you should evaluate your final trained network on the test set.

```
[54]: scores_test = two_layer_net(X_test, best_model)
print('Test accuracy: ', np.mean(np.argmax(scores_test, axis=1) == y_test))
```

```
Test accuracy: 0.536
```

```
[ ]:
```

layers

September 24, 2019

1 Modular neural nets

In the previous exercise, we computed the loss and gradient for a two-layer neural network in a single monolithic function. This isn't very difficult for a small two-layer network, but would be tedious and error-prone for larger networks. Ideally we want to build networks using a more modular design so that we can snap together different types of layers and loss functions in order to quickly experiment with different architectures.

In this exercise we will implement this approach, and develop a number of different layer types in isolation that can then be easily plugged together. For each layer we will implement **forward** and **backward** functions. The **forward** function will receive data, weights, and other parameters, and will return both an output and a **cache** object that stores data needed for the backward pass. The **backward** function will receive upstream derivatives and the cache object, and will return gradients with respect to the data and all of the weights. This will allow us to write code that looks like this:

```
def two_layer_net(X, W1, b1, W2, b2, reg):
    # Forward pass; compute scores
    s1, fc1_cache = affine_forward(X, W1, b1)
    a1, relu_cache = relu_forward(s1)
    scores, fc2_cache = affine_forward(a1, W2, b2)

    # Loss functions return data loss and gradients on scores
    data_loss, dscores = svm_loss(scores, y)

    # Compute backward pass
    da1, dW2, db2 = affine_backward(dscores, fc2_cache)
    ds1 = relu_backward(da1, relu_cache)
    dX, dW1, db1 = affine_backward(ds1, fc1_cache)

    # A real network would add regularization here

    # Return loss and gradients
    return loss, dW1, db1, dW2, db2
```

```
[67]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
```



```

from cs231n.gradient_check import eval_numerical_gradient_array, \
    eval_numerical_gradient
from cs231n.layers import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

2 Affine layer: forward

Open the file `cs231n/layers.py` and implement the `affine_forward` function.

Once you are done we will test your can test your implementation by running the following:

```

[68]: # Test the affine_forward function

num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape), \
    output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,  3.77273342]])

# Compare your output with ours. The error should be around 1e-9.

```

```
print('Testing affine_forward function:')
print('difference: ', rel_error(out, correct_out))
```

```
Testing affine_forward function:
difference: 9.769847728806635e-10
```

3 Affine layer: backward

Now implement the `affine_backward` function. You can test your implementation using numeric gradient checking.

```
[69]: # Test the affine_backward function

x = np.random.randn(10, 2, 3)
w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x,
    ↪dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w,
    ↪dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b,
    ↪dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be less than 1e-10
print('Testing affine_backward function:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing affine_backward function:
dx error: 1.389772357243565e-10
dw error: 2.0802769210228722e-09
db error: 3.283181149785737e-11
```

4 ReLU layer: forward

Implement the `relu_forward` function and test your implementation by running the following:

```
[70]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)
```

```

out, _ = relu_forward(x)
correct_out = np.array([[ 0.,          0.,          0.,          0.,          ],
                        [ 0.,          0.,          0.04545455, 0.13636364, ],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5,          ]])

# Compare your output with ours. The error should be around 1e-8
print('Testing relu_forward function:')
print('difference: ', rel_error(out, correct_out))

```

Testing relu_forward function:
 difference: 4.999999798022158e-08

5 ReLU layer: backward

Implement the `relu_backward` function and test your implementation using numeric gradient checking:

```

[71]: x = np.random.randn(10, 10)
      dout = np.random.randn(*x.shape)

      dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

      _, cache = relu_forward(x)
      dx = relu_backward(dout, cache)

      # The error should be around 1e-12
      print('Testing relu_backward function:')
      print('dx error: ', rel_error(dx_num, dx))

```

Testing relu_backward function:
 dx error: 3.2756297249955374e-12

6 Loss layers: Softmax and SVM

You implemented these loss functions in the last assignment, so we'll give them to you for free here. It's still a good idea to test them to make sure they work correctly.

```

[72]: num_classes, num_inputs = 10, 50
      x = 0.001 * np.random.randn(num_inputs, num_classes)
      y = np.random.randint(num_classes, size=num_inputs)

      dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
      loss, dx = svm_loss(x, y)

      # Test svm_loss function. Loss should be around 9 and dx error should be 1e-9
      print('Testing svm_loss:')
      print('loss: ', loss)

```

```

print('dx error: ', rel_error(dx_num, dx))

dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x,
    verbose=False)
loss, dx = softmax_loss(x, y)

# Test softmax_loss function. Loss should be 2.3 and dx error should be 1e-8
print('\nTesting softmax_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

```

```

Testing svm_loss:
loss: 8.99959980158041
dx error: 1.4021566006651672e-09

```

```

Testing softmax_loss:
loss: 2.302545580559075
dx error: 1.002114418846665e-08

```

7 Convolution layer: forward naive

We are now ready to implement the forward pass for a convolutional layer. Implement the function `conv_forward_naive` in the file `cs231n/layers.py`.

You don't have to worry too much about efficiency at this point; just write the code in whatever way you find most clear.

You can test your implementation by running the following:

```

[73]: x_shape = (2, 3, 4, 4)
      w_shape = (3, 3, 4, 4)
      x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
      w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
      b = np.linspace(-0.1, 0.2, num=3)

      conv_param = {'stride': 2, 'pad': 1}
      out, _ = conv_forward_naive(x, w, b, conv_param)
      correct_out = np.array([[[[-0.08759809, -0.10987781],
                                [-0.18387192, -0.2109216 ]],
                              [[ 0.21027089,  0.21661097],
                                [ 0.22847626,  0.23004637]],
                              [[ 0.50813986,  0.54309974],
                                [ 0.64082444,  0.67101435]]],
                             [[[-0.98053589, -1.03143541],
                                [-1.19128892, -1.24695841]],
                              [[ 0.69108355,  0.66880383],
                                [ 0.59480972,  0.56776003]],
                              [[ 2.36270298,  2.36904306],

```

```

[ 2.38090835,  2.38247847]]]]))

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))

```

```

Testing conv_forward_naive
difference:  2.2121476417505994e-08

```

8 Aside: Image processing via convolutions

As fun way to both check your implementation and gain a better understanding of the type of operation that convolutional layers can perform, we will set up an input containing two images and manually set up filters that perform common image processing operations (grayscale conversion and edge detection). The convolution forward pass will apply these operations to each of the input images. We can then visualize the results as a sanity check.

```

[74]: from scipy.misc import imread, imresize

kitten, puppy = imread('kitten.jpg'), imread('puppy.jpg')
# kitten is wide, and puppy is already square
d = kitten.shape[1] - kitten.shape[0]
kitten_cropped = kitten[:, d//2:-d//2, :]

img_size = 200 # Make this smaller if it runs too slow
x = np.zeros((2, 3, img_size, img_size))
x[0, :, :, :] = imresize(puppy, (img_size, img_size)).transpose((2, 0, 1))
x[1, :, :, :] = imresize(kitten_cropped, (img_size, img_size)).transpose((2, 0,
→1))

# Set up a convolutional weights holding 2 filters, each 3x3
w = np.zeros((2, 3, 3, 3))

# The first filter converts the image to grayscale.
# Set up the red, green, and blue channels of the filter.
w[0, 0, :, :] = [[0, 0, 0], [0, 0.3, 0], [0, 0, 0]]
w[0, 1, :, :] = [[0, 0, 0], [0, 0.6, 0], [0, 0, 0]]
w[0, 2, :, :] = [[0, 0, 0], [0, 0.1, 0], [0, 0, 0]]

# Second filter detects horizontal edges in the blue channel.
w[1, 2, :, :] = [[1, 2, 1], [0, 0, 0], [-1, -2, -1]]

# Vector of biases. We don't need any bias for the grayscale
# filter, but for the edge detection filter we want to add 128
# to each output so that nothing is negative.
b = np.array([0, 128])

```

```

# Compute the result of convolving each input in x with each filter in w,
# offsetting by b, and storing the results in out.
out, _ = conv_forward_naive(x, w, b, {'stride': 1, 'pad': 1})

def imshow_noax(img, normalize=True):
    """ Tiny helper to show images as uint8 and remove axis labels """
    if normalize:
        img_max, img_min = np.max(img), np.min(img)
        img = 255.0 * (img - img_min) / (img_max - img_min)
    plt.imshow(img.astype('uint8'))
    plt.gca().axis('off')

# Show the original images and the results of the conv operation
plt.subplot(2, 3, 1)
imshow_noax(puppy, normalize=False)
plt.title('Original image')
plt.subplot(2, 3, 2)
imshow_noax(out[0, 0])
plt.title('Grayscale')
plt.subplot(2, 3, 3)
imshow_noax(out[0, 1])
plt.title('Edges')
plt.subplot(2, 3, 4)
imshow_noax(kitten_cropped, normalize=False)
plt.subplot(2, 3, 5)
imshow_noax(out[1, 0])
plt.subplot(2, 3, 6)
imshow_noax(out[1, 1])
plt.show()

```

```

↳ -----
ImportError                                Traceback (most recent call↳
↳ last)

<ipython-input-74-92a8b739741c> in <module>
----> 1 from scipy.misc import imread, imresize
      2
      3 kitten, puppy = imread('kitten.jpg'), imread('puppy.jpg')
      4 # kitten is wide, and puppy is already square
      5 d = kitten.shape[1] - kitten.shape[0]

ImportError: cannot import name 'imread'

```

9 Convolution layer: backward naive

Next you need to implement the function `conv_backward_naive` in the file `cs231n/layers.py`. As usual, we will check your implementation with numeric gradient checking.

```
[78]: x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b,
    ↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b,
    ↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b,
    ↪conv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9'
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))
```

Testing conv_backward_naive function

dx error: 2.6599967595550446e-09

dw error: 2.2541074080896546e-10

db error: 9.5467474823066e-12

10 Max pooling layer: forward naive

The last layer we need for a basic convolutional neural network is the max pooling layer. First implement the forward pass in the function `max_pool_forward_naive` in the file `cs231n/layers.py`.

```
[79]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]]],
```

```

[[[-0.02736842, -0.01263158],
  [ 0.03157895,  0.04631579]]],
[[[ 0.09052632,  0.10526316],
  [ 0.14947368,  0.16421053]],
[[ 0.20842105,  0.22315789],
  [ 0.26736842,  0.28210526]],
[[ 0.32631579,  0.34105263],
  [ 0.38526316,  0.4         ]]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

```

Testing max_pool_forward_naive function:
 difference: 4.1666665157267834e-08

11 Max pooling layer: backward naive

Implement the backward pass for a max pooling layer in the function `max_pool_backward_naive` in the file `cs231n/layers.py`. As always we check the correctness of the backward pass using numerical gradient checking.

```

[80]: x = np.random.randn(3, 2, 8, 8)
      dout = np.random.randn(3, 2, 4, 4)
      pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

      dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x,
      ↪pool_param)[0], x, dout)

      out, cache = max_pool_forward_naive(x, pool_param)
      dx = max_pool_backward_naive(dout, cache)

      # Your error should be around 1e-12
      print('Testing max_pool_backward_naive function:')
      print('dx error: ', rel_error(dx, dx_num))

```

Testing max_pool_backward_naive function:
 dx error: 3.275642258527761e-12

12 Fast layers

Making convolution and pooling layers fast can be challenging. To spare you the pain, we've provided fast implementations of the forward and backward passes for convolution and pooling layers in the file `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:


```
python setup.py build_ext --inplace
```

The API for the fast versions of the convolution and pooling layers is exactly the same as the naive versions that you implemented above: the forward pass receives data, weights, and parameters and produces outputs and a cache object; the backward pass receives upstream derivatives and the cache object and produces gradients with respect to the data and weights.

NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the following:

```
[81]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
      from time import time

      x = np.random.randn(100, 3, 31, 31)
      w = np.random.randn(25, 3, 3, 3)
      b = np.random.randn(25,)
      dout = np.random.randn(100, 25, 16, 16)
      conv_param = {'stride': 2, 'pad': 1}

      t0 = time()
      out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
      t1 = time()
      out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
      t2 = time()

      print('Testing conv_forward_fast:')
      print('Naive: %fs' % (t1 - t0))
      print('Fast: %fs' % (t2 - t1))
      print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
      print('Difference: ', rel_error(out_naive, out_fast))

      t0 = time()
      dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
      t1 = time()
      dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
      t2 = time()

      print('\nTesting conv_backward_fast:')
      print('Naive: %fs' % (t1 - t0))
      print('Fast: %fs' % (t2 - t1))
      print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
      print('dx difference: ', rel_error(dx_naive, dx_fast))
      print('dw difference: ', rel_error(dw_naive, dw_fast))
      print('db difference: ', rel_error(db_naive, db_fast))
```

Testing conv_forward_fast:
Naive: 4.739521s
Fast: 0.008533s
Speedup: 555.434255x
Difference: 7.053358279005642e-11

Testing conv_backward_fast:
Naive: 6.787412s
Fast: 0.011820s
Speedup: 574.238946x
dx difference: 5.415099860467299e-12
dw difference: 1.3570073863477166e-12
db difference: 2.3571176246445243e-14

```
[82]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
```

Testing pool_forward_fast:
Naive: 0.354986s
fast: 0.001795s
speedup: 197.757870x
difference: 0.0

```
Testing pool_backward_fast:
Naive: 0.366232s
speedup: 39.574572x
dx difference: 0.0
```

13 Sandwich layers

There are a couple common layer “sandwiches” that frequently appear in ConvNets. For example convolutional layers are frequently followed by ReLU and pooling, and affine layers are frequently followed by ReLU. To make it more convenient to use these common patterns, we have defined several convenience layers in the file `cs231n/layer_utils.py`. Lets grad-check them to make sure that they work correctly:

```
[83]: from cs231n.layer_utils import conv_relu_pool_forward, conv_relu_pool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w,
    ↪b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w,
    ↪b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w,
    ↪b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool_forward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing conv_relu_pool_forward:
dx error: 6.702093662796137e-09
dw error: 1.5446117758365741e-09
db error: 2.0456884116571853e-11
```

```
[84]: from cs231n.layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
```

```

b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b,
    ↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b,
    ↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b,
    ↪conv_param)[0], b, dout)

print('Testing conv_relu_forward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

Testing conv_relu_forward:

dx error: 2.0945883679826754e-08

dw error: 4.214221674271771e-09

db error: 6.630038270484846e-12

[85]: `from cs231n.layer_utils import affine_relu_forward, affine_relu_backward`

```

x = np.random.randn(2, 3, 4)
w = np.random.randn(12, 10)
b = np.random.randn(10)
dout = np.random.randn(2, 10)

out, cache = affine_relu_forward(x, w, b)
dx, dw, db = affine_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w,
    ↪b)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w,
    ↪b)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w,
    ↪b)[0], b, dout)

print('Testing affine_relu_forward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

Testing affine_relu_forward:

```
dx error: 3.141513976758779e-10  
dw error: 6.355226245075724e-10  
db error: 3.2755971928120225e-12
```

```
[ ]:
```

convnet

September 24, 2019

1 Train a ConvNet!

We now have a generic solver and a bunch of modularized layers. It's time to put it all together, and train a ConvNet to recognize the classes in CIFAR-10. In this notebook we will walk you through training a simple two-layer ConvNet and then set you free to build the best net that you can to perform well on CIFAR-10.

Open up the file `cs231n/classifiers/convnet.py`; you will see that the `two_layer_convnet` function computes the loss and gradients for a two-layer ConvNet. Note that this function uses the “sandwich” layers defined in `cs231n/layer_utils.py`.

```
[1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from cs231n.classifier_trainer import ClassifierTrainer
from cs231n.gradient_check import eval_numerical_gradient
from cs231n.classifiers.convnet import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
[2]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
```

Load the CIFAR-10 dataset from disk and perform preprocessing to prepare it for the two-layer neural net classifier. These are the same steps as we used for the SVM, but condensed to a single function.

"""

Load the raw CIFAR-10 data

cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

Subsample the data

mask = range(num_training, num_training + num_validation)

X_val = X_train[mask]

y_val = y_train[mask]

mask = range(num_training)

X_train = X_train[mask]

y_train = y_train[mask]

mask = range(num_test)

X_test = X_test[mask]

y_test = y_test[mask]

Normalize the data: subtract the mean image

mean_image = np.mean(X_train, axis=0)

X_train -= mean_image

X_val -= mean_image

X_test -= mean_image

Transpose so that channels come first

X_train = X_train.transpose(0, 3, 1, 2).copy()

X_val = X_val.transpose(0, 3, 1, 2).copy()

x_test = X_test.transpose(0, 3, 1, 2).copy()

return X_train, y_train, X_val, y_val, X_test, y_test

Invoke the above function to get our data.

X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()

print('Train data shape: ', X_train.shape)

print('Train labels shape: ', y_train.shape)

print('Validation data shape: ', X_val.shape)

print('Validation labels shape: ', y_val.shape)

print('Test data shape: ', X_test.shape)

print('Test labels shape: ', y_test.shape)

Train data shape: (49000, 3, 32, 32)

Train labels shape: (49000,)

Validation data shape: (1000, 3, 32, 32)

Validation labels shape: (1000,)

Test data shape: (1000, 32, 32, 3)

Test labels shape: (1000,)

2 Sanity check loss

After you build a new network, one of the first things you should do is sanity check the loss. When we use the softmax loss, we expect the loss for random weights (and no regularization) to be about $\log(C)$ for C classes. When we add regularization this should go up.

```
[3]: model = init_two_layer_convnet()

X = np.random.randn(100, 3, 32, 32)
y = np.random.randint(10, size=100)

loss, _ = two_layer_convnet(X, model, y, reg=0)

# Sanity check: Loss should be about log(10) = 2.3026
print('Sanity check loss (no regularization): ', loss)

# Sanity check: Loss should go up when you add regularization
loss, _ = two_layer_convnet(X, model, y, reg=1)
print('Sanity check loss (with regularization): ', loss)
```

Sanity check loss (no regularization): 2.302678587918066

Sanity check loss (with regularization): 2.3447926824948357

3 Gradient check

After the loss looks reasonable, you should always use numeric gradient checking to make sure that your backward pass is correct. When you use numeric gradient checking you should use a small amount of artificial data and a small number of neurons at each layer.

```
[4]: num_inputs = 2
input_shape = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_shape)
y = np.random.randint(num_classes, size=num_inputs)

model = init_two_layer_convnet(num_filters=3, filter_size=3,
    ↪input_shape=input_shape)
loss, grads = two_layer_convnet(X, model, y)
for param_name in sorted(grads):
    f = lambda _: two_layer_convnet(X, model, y)[0]
    param_grad_num = eval_numerical_gradient(f, model[param_name],
    ↪verbose=False, h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
```



```
print('%s max relative error: %e' % (param_name, rel_error(param_grad_num, ↵
↵grads[param_name])))
```

```
W1 max relative error: 5.290520e-06
W2 max relative error: 1.848474e-05
b1 max relative error: 5.200847e-08
b2 max relative error: 6.433921e-10
```

4 Overfit small data

A nice trick is to train your model with just a few training samples. You should be able to overfit small datasets, which will result in very high training accuracy and comparatively low validation accuracy.

```
[5]: # Use a two-layer ConvNet to overfit 50 training examples.
```

```
model = init_two_layer_convnet()
trainer = ClassifierTrainer()
best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
    X_train[:50], y_train[:50], X_val, y_val, model, two_layer_convnet,
    reg=0.001, momentum=0.9, learning_rate=0.0001, batch_size=10, ↵
    ↵num_epochs=10,
    verbose=True)
```

```
starting iteration 0
Finished epoch 0 / 10: cost 2.292191, train: 0.120000, val 0.118000, lr
1.000000e-04
Finished epoch 1 / 10: cost 2.201300, train: 0.260000, val 0.098000, lr
9.500000e-05
Finished epoch 2 / 10: cost 1.918935, train: 0.280000, val 0.096000, lr
9.025000e-05
starting iteration 10
Finished epoch 3 / 10: cost 1.232655, train: 0.400000, val 0.166000, lr
8.573750e-05
Finished epoch 4 / 10: cost 1.088217, train: 0.560000, val 0.156000, lr
8.145062e-05
starting iteration 20
Finished epoch 5 / 10: cost 0.828684, train: 0.680000, val 0.189000, lr
7.737809e-05
Finished epoch 6 / 10: cost 1.657847, train: 0.740000, val 0.166000, lr
7.350919e-05
starting iteration 30
Finished epoch 7 / 10: cost 0.608962, train: 0.860000, val 0.177000, lr
6.983373e-05
Finished epoch 8 / 10: cost 0.350635, train: 0.900000, val 0.146000, lr
6.634204e-05
starting iteration 40
```

```

Finished epoch 9 / 10: cost 0.259157, train: 0.900000, val 0.167000, lr
6.302494e-05
Finished epoch 10 / 10: cost 0.797573, train: 0.880000, val 0.161000, lr
5.987369e-05
finished optimization. best validation accuracy: 0.189000

```

Plotting the loss, training accuracy, and validation accuracy should show clear overfitting:

```

[ ]: plt.subplot(2, 1, 1)
      plt.plot(loss_history)
      plt.xlabel('iteration')
      plt.ylabel('loss')

      plt.subplot(2, 1, 2)
      plt.plot(train_acc_history)
      plt.plot(val_acc_history)
      plt.legend(['train', 'val'], loc='upper left')
      plt.xlabel('epoch')
      plt.ylabel('accuracy')
      plt.show()

```

5 Train the net

Once the above works, training the net is the next thing to try. You can set the `acc_frequency` parameter to change the frequency at which the training and validation set accuracies are tested. If your parameters are set properly, you should see the training and validation accuracy start to improve within a hundred iterations, and you should be able to train a reasonable model with just one epoch.

Using the parameters below you should be able to get around 50% accuracy on the validation set.

```

[6]: model = init_two_layer_convnet(filter_size=7)
      trainer = ClassifierTrainer()
      best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
          X_train, y_train, X_val, y_val, model, two_layer_convnet,
          reg=0.001, momentum=0.9, learning_rate=0.0001, batch_size=50,
          ↪num_epochs=1,
          acc_frequency=50, verbose=True)

```

```

starting iteration 0
Finished epoch 0 / 1: cost 2.326596, train: 0.115000, val 0.101000, lr
1.000000e-04
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
Finished epoch 0 / 1: cost 1.722926, train: 0.341000, val 0.349000, lr

```

1.000000e-04
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
starting iteration 100
Finished epoch 0 / 1: cost 1.792310, train: 0.334000, val 0.338000, lr
1.000000e-04
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
Finished epoch 0 / 1: cost 1.840112, train: 0.395000, val 0.363000, lr
1.000000e-04
starting iteration 160
starting iteration 170
starting iteration 180
starting iteration 190
starting iteration 200
Finished epoch 0 / 1: cost 1.764864, train: 0.379000, val 0.410000, lr
1.000000e-04
starting iteration 210
starting iteration 220
starting iteration 230
starting iteration 240
starting iteration 250
Finished epoch 0 / 1: cost 1.769096, train: 0.427000, val 0.419000, lr
1.000000e-04
starting iteration 260
starting iteration 270
starting iteration 280
starting iteration 290
starting iteration 300
Finished epoch 0 / 1: cost 1.524768, train: 0.435000, val 0.453000, lr
1.000000e-04
starting iteration 310
starting iteration 320
starting iteration 330
starting iteration 340
starting iteration 350
Finished epoch 0 / 1: cost 2.337207, train: 0.399000, val 0.367000, lr
1.000000e-04
starting iteration 360
starting iteration 370
starting iteration 380
starting iteration 390
starting iteration 400

Finished epoch 0 / 1: cost 1.473121, train: 0.465000, val 0.458000, lr 1.000000e-04
starting iteration 410
starting iteration 420
starting iteration 430
starting iteration 440
starting iteration 450
Finished epoch 0 / 1: cost 1.717805, train: 0.414000, val 0.457000, lr 1.000000e-04
starting iteration 460
starting iteration 470
starting iteration 480
starting iteration 490
starting iteration 500
Finished epoch 0 / 1: cost 1.587554, train: 0.513000, val 0.491000, lr 1.000000e-04
starting iteration 510
starting iteration 520
starting iteration 530
starting iteration 540
starting iteration 550
Finished epoch 0 / 1: cost 1.648589, train: 0.468000, val 0.452000, lr 1.000000e-04
starting iteration 560
starting iteration 570
starting iteration 580
starting iteration 590
starting iteration 600
Finished epoch 0 / 1: cost 1.650349, train: 0.415000, val 0.429000, lr 1.000000e-04
starting iteration 610
starting iteration 620
starting iteration 630
starting iteration 640
starting iteration 650
Finished epoch 0 / 1: cost 1.746694, train: 0.470000, val 0.497000, lr 1.000000e-04
starting iteration 660
starting iteration 670
starting iteration 680
starting iteration 690
starting iteration 700
Finished epoch 0 / 1: cost 1.624673, train: 0.495000, val 0.507000, lr 1.000000e-04
starting iteration 710
starting iteration 720
starting iteration 730
starting iteration 740

```

starting iteration 750
Finished epoch 0 / 1: cost 1.218479, train: 0.533000, val 0.478000, lr
1.000000e-04
starting iteration 760
starting iteration 770
starting iteration 780
starting iteration 790
starting iteration 800
Finished epoch 0 / 1: cost 1.675719, train: 0.486000, val 0.463000, lr
1.000000e-04
starting iteration 810
starting iteration 820
starting iteration 830
starting iteration 840
starting iteration 850
Finished epoch 0 / 1: cost 1.413279, train: 0.468000, val 0.486000, lr
1.000000e-04
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
Finished epoch 0 / 1: cost 1.548007, train: 0.471000, val 0.453000, lr
1.000000e-04
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950
Finished epoch 0 / 1: cost 2.070196, train: 0.524000, val 0.469000, lr
1.000000e-04
starting iteration 960
starting iteration 970
Finished epoch 1 / 1: cost 1.489284, train: 0.429000, val 0.433000, lr
9.500000e-05
finished optimization. best validation accuracy: 0.507000

```

6 Visualize weights

We can visualize the convolutional weights from the first layer. If everything worked properly, these will usually be edges and blobs of various colors and orientations.


```

[7]: from cs231n.vis_utils import visualize_grid

grid = visualize_grid(best_model['W1'].transpose(0, 2, 3, 1))
plt.imshow(grid.astype('uint8'))

```

```
[7]: <matplotlib.image.AxesImage at 0x11ba67e48>
```



```
[ ]:
```

softmax-classifier

September 26, 2019

0.1 PyTorch data

PyTorch comes with a nice paradigm for dealing with data which we'll use here. A PyTorch [Dataset](#) knows where to find data in its raw form (files on disk) and how to load individual examples into Python datastructures. A PyTorch [DataLoader](#) takes a dataset and offers a variety of ways to sample batches from that dataset.

Take a moment to browse through the CIFAR10 Dataset in `2_pytorch/cifar10.py`, read the [DataLoader](#) documentation linked above, and see how these are used in the section of `train.py` that loads data. Note that in the first part of the homework we subtracted a mean CIFAR10 image from every image before feeding it in to our models. Here we subtract a constant color instead. Both methods are seen in practice and work equally well.

PyTorch provides lots of vision datasets which can be imported directly from [torchvision.datasets](#). Also see [torchtext](#) for natural language datasets.

0.2 Softmax Classifier in PyTorch

In PyTorch Deep Learning building blocks are implemented in the neural network module [torch.nn](#) (usually imported as `nn`). A PyTorch model is typically a subclass of [nn.Module](#) and thereby gains a multitude of features. Because your logistic regressor is an `nn.Module` all of its parameters and sub-modules are accessible through the `.parameters()` and `.modules()` methods.

Now implement a softmax classifier by filling in the marked sections of `models/softmax.py`.

The main driver for this question is `train.py`. It reads arguments and model hyperparameter from the command line, loads CIFAR10 data and the specified model (in this case, softmax). Using the optimizer initialized with appropriate hyperparameters, it trains the model and reports performance on test data.

Complete the following couple of sections in `train.py`: 1. Initialize an optimizer from the `torch.optim` package 2. Update the parameters in model using the optimizer initialized above

At this point all of the components required to train the softmax classifier are complete for the softmax classifier. Now run

```
$ run_softmax.sh
```

to train a model and save it to `softmax.pt`. This will also produce a `softmax.log` file which contains training details which we will visualize below.

Note: You may want to adjust the hyperparameters specified in `run_softmax.sh` to get reasonable performance.

0.3 Visualizing the PyTorch model

```
[6]: # Assuming that you have completed training the classifier, let us plot the
      ↪ training loss vs. iteration. This is an
      # example to show a simple way to log and plot data from PyTorch.

      # we need matplotlib to plot the graphs for us!
      import matplotlib
      # This is needed to save images
      matplotlib.use('Agg')
      import matplotlib.pyplot as plt
      %matplotlib inline
```

//anaconda3/envs/cs7643/lib/python3.6/site-packages/ipykernel_launcher.py:7:

UserWarning:

This call to matplotlib.use() has no effect because the backend has already been chosen; matplotlib.use() must be called *before* pylab, matplotlib.pyplot, or matplotlib.backends is imported for the first time.

The backend was *originally* set to 'module://ipykernel.pylab.backend_inline' by the following code:

```
File "//anaconda3/envs/cs7643/lib/python3.6/runpy.py", line 193, in
_run_module_as_main
    "__main__", mod_spec)
File "//anaconda3/envs/cs7643/lib/python3.6/runpy.py", line 85, in _run_code
    exec(code, run_globals)
File "//anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel_launcher.py", line 16, in <module>
    app.launch_new_instance()
File "//anaconda3/envs/cs7643/lib/python3.6/site-
packages/traitlets/config/application.py", line 658, in launch_instance
    app.start()
File "//anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel/kernelapp.py", line 563, in start
    self.io_loop.start()
File "//anaconda3/envs/cs7643/lib/python3.6/site-
packages/tornado/platform/asyncio.py", line 148, in start
    self.asyncio_loop.run_forever()
File "//anaconda3/envs/cs7643/lib/python3.6/asyncio/base_events.py", line 421,
in run_forever
    self._run_once()
File "//anaconda3/envs/cs7643/lib/python3.6/asyncio/base_events.py", line
1425, in _run_once
    handle._run()
File "//anaconda3/envs/cs7643/lib/python3.6/asyncio/events.py", line 127, in
_run
    self._callback(*self._args)
File "//anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/ioloop.py",
```



```

line 690, in <lambda>
    lambda f: self._run_callback(functools.partial(callback, future))
File "///anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/ioloop.py",
line 743, in _run_callback
    ret = callback()
File "///anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/gen.py",
line 787, in inner
    self.run()
File "///anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/gen.py",
line 748, in run
    yielded = self.gen.send(value)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel/kernelbase.py", line 365, in process_one
    yield gen.maybe_future(dispatch(*args))
File "///anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/gen.py",
line 209, in wrapper
    yielded = next(result)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel/kernelbase.py", line 272, in dispatch_shell
    yield gen.maybe_future(handler(stream, idents, msg))
File "///anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/gen.py",
line 209, in wrapper
    yielded = next(result)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel/kernelbase.py", line 542, in execute_request
    user_expressions, allow_stdin,
File "///anaconda3/envs/cs7643/lib/python3.6/site-packages/tornado/gen.py",
line 209, in wrapper
    yielded = next(result)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel/ipkernel.py", line 294, in do_execute
    res = shell.run_cell(code, store_history=store_history, silent=silent)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/ipykernel/zmqshell.py", line 536, in run_cell
    return super(ZMQInteractiveShell, self).run_cell(*args, **kwargs)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/interactiveshell.py", line 2855, in run_cell
    raw_cell, store_history, silent, shell_futures)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/interactiveshell.py", line 2881, in _run_cell
    return runner(coro)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/async_helpers.py", line 68, in _pseudo_sync_runner
    coro.send(None)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/interactiveshell.py", line 3058, in run_cell_async
    interactivity=interactivity, compiler=compiler, result=result)
File "///anaconda3/envs/cs7643/lib/python3.6/site-

```

```

packages/IPython/core/interactiveshell.py", line 3249, in run_ast_nodes
    if (await self.run_code(code, result,  async_=asy)):
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/interactiveshell.py", line 3326, in run_code
    exec(code_obj, self.user_global_ns, self.user_ns)
File "<ipython-input-4-b4795f515c11>", line 9, in <module>
    get_ipython().run_line_magic('matplotlib', 'inline')
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/interactiveshell.py", line 2314, in run_line_magic
    result = fn(*args, **kwargs)
File "<///anaconda3/envs/cs7643/lib/python3.6/site-
packages/decorator.py:decorator-gen-108>", line 2, in matplotlib
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/magic.py", line 187, in <lambda>
    call = lambda f, *a, **k: f(*a, **k)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/magics/pylab.py", line 99, in matplotlib
    gui, backend = self.shell.enable_matplotlib(args.gui.lower() if
isinstance(args.gui, str) else args.gui)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/interactiveshell.py", line 3414, in enable_matplotlib
    pt.activate_matplotlib(backend)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/IPython/core/pylabtools.py", line 314, in activate_matplotlib
    matplotlib.pyplot.switch_backend(backend)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/matplotlib/pyplot.py", line 231, in switch_backend
    matplotlib.use(newbackend, warn=False, force=True)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/matplotlib/__init__.py", line 1410, in use
    reload(sys.modules['matplotlib.backends'])
File "///anaconda3/envs/cs7643/lib/python3.6/importlib/__init__.py", line 166,
in reload
    _bootstrap._exec(spec, module)
File "///anaconda3/envs/cs7643/lib/python3.6/site-
packages/matplotlib/backends/__init__.py", line 16, in <module>
    line for line in traceback.format_stack()

```

```
import sys
```

```

[8]: # Parse the train and val losses one line at a time.
import re
# regexes to find train and val losses on a line
float_regex = r'[-+]?(\d+(\.\d*)?|\.\d+)([eE][-+]?\d+)?'
train_loss_re = re.compile('.*Train Loss: ({})'.format(float_regex))
val_loss_re = re.compile('.*Val Loss: ({})'.format(float_regex))

```

```

val_acc_re = re.compile('.*Val Acc: ({})*'.format(float_regex))
# extract one loss for each logged iteration
train_losses = []
val_losses = []
val_accs = []
# NOTE: You may need to change this file name.
with open('softmax.log', 'r') as f:
    for line in f:
        train_match = train_loss_re.match(line)
        val_match = val_loss_re.match(line)
        val_acc_match = val_acc_re.match(line)
        if train_match:
            train_losses.append(float(train_match.group(1)))
        if val_match:
            val_losses.append(float(val_match.group(1)))
        if val_acc_match:
            val_accs.append(float(val_acc_match.group(1)))



```

```

[9]: fig = plt.figure()
plt.plot(train_losses, label='train')
plt.plot(val_losses, label='val')
plt.title('Softmax Learning Curve')
plt.ylabel('loss')
plt.legend()
fig.savefig('softmax_lossvstrain.png')




fig = plt.figure()
plt.plot(val_accs, label='val')
plt.title('Softmax Validation Accuracy During Training')
plt.ylabel('accuracy')
plt.legend()
fig.savefig('softmax_valaccuracy.png')

```


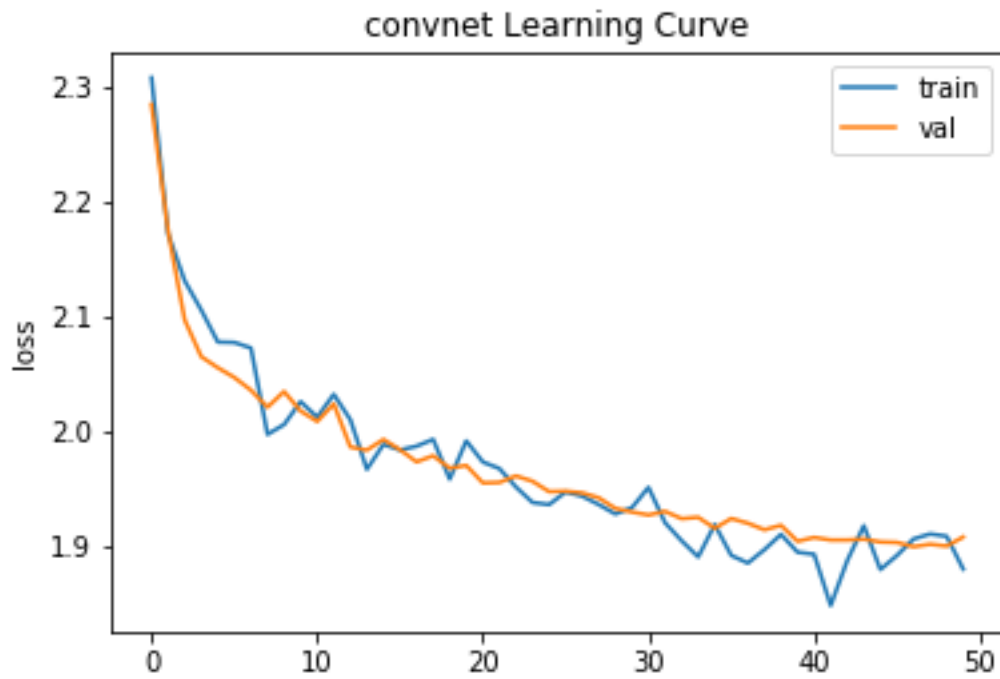


[:

Q8.6 Two-Layer



Q8.7 ConvNet



Q8.8 Experiment

